

# CS 4495 Computer Vision

## *Linear Filtering 2: Templates, Edges*

---

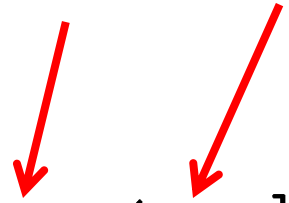
**Aaron Bobick**

School of Interactive  
Computing



# Last time: Convolution


- Convolution:
  - Flip the filter in both dimensions (right to left, bottom to top)
  - Then apply cross-correlation

$$G[i, j] = \sum_{u=-k}^k \sum_{v=-k}^k H[u, v] F[i - u, j - v]$$


$$G = H * F$$



*Notation for  
convolution  
operator*



# Convolution vs. correlation

## Convolution

$$G[i, j] = \sum_{u=-k}^k \sum_{v=-k}^k H[u, v] F[i - u, j - v]$$

$$G = H * F$$

## (Cross-)correlation

$$G[i, j] = \sum_{u=-k}^k \sum_{v=-k}^k H[u, v] F[i + u, j + v]$$

$$G = H \otimes F$$



- When  $H$  is symmetric, no difference. We tend to use the terms interchangeably.
- Convolution with an impulse (centered at 0,0) is the identity

# Filters for features


- Previously, thinking of filtering as a way to remove or reduce **noise**
- Now, consider how filters will allow us to abstract higher-level “**features**”.
  - Map raw pixels to an intermediate representation that will be used for subsequent processing
  - Goal: reduce amount of data, discard redundancy, preserve what's useful



# Template matching


- Filters as **templates**:

Note that filters look like the effects they are intended to find --- “matched filters”




- Use (*normalized*) cross-correlation score to find a given pattern (template) in the image.
  - Normalization needed to control for relative brightness. More in problem sets.

# Template matching




**Scene**




**Template (mask)**

**A toy example**

# Template matching




**Detected template**




**Template**

# Template matching



**Detected template**



**Correlation map**



# Where's Waldo?



**Scene**



**Template**



# Where's Waldo?



**Template**

**Detected template**

# Template demo...


- In directory C:\Bobick\matlab\CS4495\Filter
- echodemo waldotemplate



# Where's Waldo?



**Detected template**



**Correlation map**

# Template matching




**Scene**



**Template**

What if the template is not identical to some subimage in the scene?

# Template matching



**Detected template**






**Template**

Match can be meaningful, if scale, orientation, and general appearance is right.

# Generic features...


- When looking for a specific object or pattern, the **features** can be defined for that pattern – we will do this later in the course for specific object recognition.
- But for **generic** images, what would be good features? What are the parts or properties of the image that encode its “meaning” for human (or other biological) observers?
- Some examples of greatly reduced images...



*Edges seem to be important...*



# Origin of Edges




- Edges are caused by a variety of factors
- Information theory view: edges encode change, change is what is hard to predict, therefore edges efficiently encode an image

# In a real image

Reflectance change:  
appearance  
information, texture


Change in surface  
orientation: shape





Depth discontinuity:  
object boundary

Cast shadows

# Contrast and invariance





# Edge detection




- Convert a 2D image into a set of curves
  - Extracts salient features of the scene
  - More compact than pixels

# Edge detection



- How can you tell that a pixel is on an edge?

# Images as functions...



- Edges look like steep cliffs

# Edge Detection

Basic idea: look for a neighborhood with strong signs of change.

Problems:

- neighborhood size
- how to detect change


81	82	26	24
82	33	25	25
81	82	26	24




# Derivatives and edges

An edge is a place of rapid change in the image intensity function.


image



intensity function  
(along horizontal scanline)



first derivative



edges correspond to  
extrema of derivative






# Differential Operators

- Differential operators – here we mean some operation that when applied to the image returns some derivatives.
- We will model these “operators” as masks/kernels which when applied to the image yields a new function that is the image *gradient function*.
- We will then threshold the this *gradient function* to select the edge pixels.
- Which brings us to the question:  
*What's a gradient?*

# Image gradient

The gradient of an image:

$$\nabla f = \left[ \frac{\partial f}{\partial x}, \frac{\partial f}{\partial y} \right]$$



The gradient points in the direction of most rapid increase in intensity

The gradient direction is given by:

$$\theta = \tan^{-1} \left( \frac{\partial f}{\partial y} / \frac{\partial f}{\partial x} \right)$$

- how does this relate to the direction of the edge?

The *edge strength* is given by the gradient magnitude

$$\|\nabla f\| = \sqrt{\left( \frac{\partial f}{\partial x} \right)^2 + \left( \frac{\partial f}{\partial y} \right)^2}$$

# Discrete gradient

For 2D function,  $f(x,y)$ , the partial derivative is:

$$\frac{\partial f(x, y)}{\partial x} = \lim_{\varepsilon \rightarrow 0} \frac{f(x + \varepsilon, y) - f(x, y)}{\varepsilon}$$


For discrete data, we can approximate using *finite* differences:

$$\frac{\partial f(x, y)}{\partial x} \approx \frac{f(x + 1, y) - f(x, y)}{1}$$

$$\approx f(x + 1, y) - f(x, y) \quad \text{“right derivative”}$$

*But is it???*

# Finite differences




# Partial derivatives of an image





$$\frac{\partial f(x, y)}{\partial x}$$

$$\partial x$$



$$\frac{\partial f(x, y)}{\partial y}$$

$$\partial y$$



Which shows changes with respect to x?  
(showing correlation filters)

# Differentiation and convolution

- For 2D function,  $f(x,y)$ , the partial derivative is:

$$\frac{\partial f(x, y)}{\partial x} = \lim_{\varepsilon \rightarrow 0} \frac{f(x + \varepsilon, y) - f(x, y)}{\varepsilon}$$

- For discrete data, we can approximate using finite differences:

$$\frac{\partial f(x, y)}{\partial x} \approx \frac{f(x + 1, y) - f(x, y)}{1}$$

- To implement above as convolution, what would be the associated filter?

# The discrete gradient

- We want an “operator” (mask/kernel) that we can apply to the image that implements:

$$\frac{\partial f(x, y)}{\partial x} = \lim_{\varepsilon \rightarrow 0} \frac{f(x + \varepsilon, y) - f(x, y)}{\varepsilon}$$

How would you implement this as a cross-correlation?  
(not flipped)

0	0
-1	+1
0	0

$H$

*Not symmetric  
around image  
point; which is  
“middle” pixel?*

0	0	0
-1/2	0	+1/2
0	0	0

$H$

*Average of  
“left” and  
“right”  
derivative .  
See?*

# Example: Sobel operator

$$\frac{1}{8} \begin{bmatrix} -1 & 0 & 1 \\ -2 & 0 & 2 \\ -1 & 0 & 1 \end{bmatrix}$$

$S_x$

$$\frac{1}{8} \begin{bmatrix} 1 & 2 & 1 \\ 0 & 0 & 0 \\ -1 & -2 & -1 \end{bmatrix}$$

$S_y$

On a pixel of the image  $I$

- Let  $g_x$  be the response to mask  $S_x$  (sometimes  $* 1/8$ )
- Let  $g_y$  be the response to mask  $S_y$

What is the gradient?

(Sobel) Gradient is  $\nabla I = [g_x \ g_y]^T$

$$g = (g_x^2 + g_y^2)^{1/2}$$


is the gradient magnitude.

$$\theta = \text{atan2}(g_y, g_x)$$


is the gradient direction.




# Sobel Operator on Blocks Image



original image



gradient  
magnitude



thresholded  
gradient  
magnitude

# Some Well-Known Masks for Computing Gradients

- Sobel:

$S_x$		
-1	0	1
-2	0	2
-1	0	1

$S_y$		
1	2	1
0	0	0
-1	-2	-1
- Prewitt:

-1	0	1
-1	0	1
-1	0	1


1	1	1
0	0	0
-1	-1	-1
- Roberts

0	1
-1	0

1	0
0	-1


# Matlab does edges

```
>> My = fspecial('sobel');  
>> outim = imfilter(double(im), My);  
>> imagesc(outim);  
>> colormap gray;
```




# But...

- Consider a single row or column of the image
  - Plotting intensity as a function of  $x$




- Apply derivative operator....




Uh, where's  
the edge?

# Finite differences responding to noise



Increasing noise ->  
(this is zero mean additive gaussian noise)

# Solution: smooth first




Where is the edge? Look for peaks in  $\frac{\partial}{\partial x}(h * f)$

# Derivative theorem of convolution


- This saves us one operation:  $\frac{\partial}{\partial x} (h * f) = \left( \frac{\partial}{\partial x} h \right) * f$

Sigma = 50


$f$




$h$



$\frac{\partial}{\partial x} h$



$\left( \frac{\partial}{\partial x} h \right) * f$





How can we find (local) maxima of a function?


# 2<sup>nd</sup> derivative of Gaussian

- Consider  $\frac{\partial^2}{\partial x^2} (h * f)$


$f$



$\frac{\partial^2}{\partial x^2} h$



$\frac{\partial^2}{\partial x^2} (h * f)$



Where is the edge?      Zero-crossings of bottom graph




# What about 2D?

# Derivative of Gaussian filter


$$(I \otimes g) \otimes h = I \otimes (g \otimes h)$$

$$\begin{bmatrix} 0.0030 & 0.0133 & 0.0219 & 0.0133 & 0.0030 \\ 0.0133 & 0.0596 & 0.0983 & 0.0596 & 0.0133 \\ 0.0219 & 0.0983 & 0.1621 & 0.0983 & 0.0219 \\ 0.0133 & 0.0596 & 0.0983 & 0.0596 & 0.0133 \\ 0.0030 & 0.0133 & 0.0219 & 0.0133 & 0.0030 \end{bmatrix} \otimes \begin{bmatrix} 1 & -1 \end{bmatrix}$$




Why is this preferable?


# Derivative of Gaussian filters




x-direction



*Is this for correlation  
or convolution?*



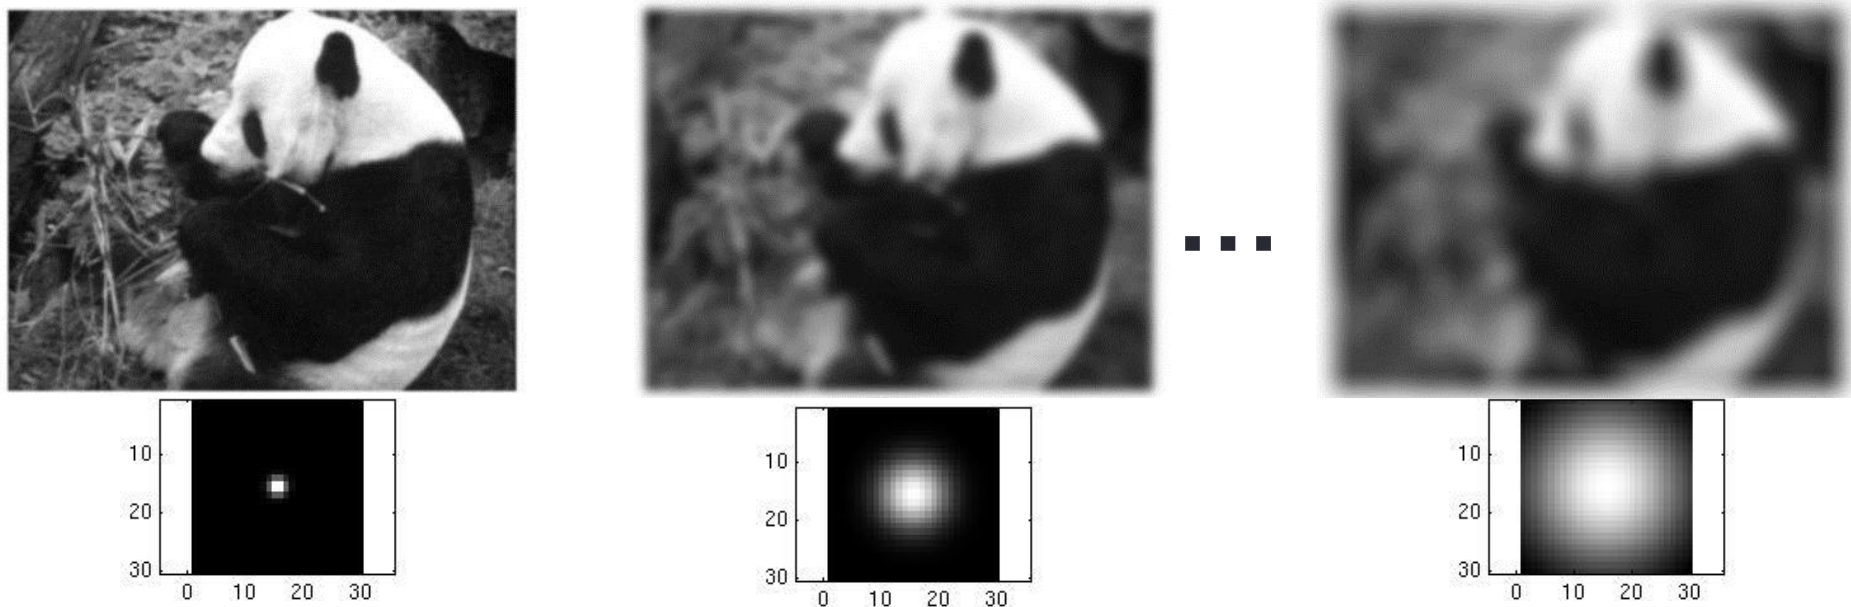
y-direction



*And for y it's always  
a problem!*

# Smoothing with a Gaussian



Parameter  $\sigma$  is the “scale” / “width” / “spread” of the Gaussian kernel, and controls the amount of smoothing.




```
for sigma=1:3:10
    h = fspecial('gaussian', fsize, sigma);
    out = imfilter(im, h);
    imshow(out);
    pause;
end
```

# Smoothing with a Gaussian


Recall: parameter  $\sigma$  is the “scale” / “width” / “spread” of the Gaussian kernel, and controls the amount of smoothing.




...



# Effect of $\sigma$ on derivatives



$\sigma = 1$  pixel



$\sigma = 3$  pixels

The apparent structures differ depending on Gaussian's scale parameter.

Larger values: larger scale edges detected  
Smaller values: finer features detected



# Gradients -> edges

- Primary edge detection steps:
  - 1. Smoothing: suppress noise
  - 2. Edge “enhancement”: filter for contrast
  - 3. Edge localization
    - Determine which local maxima from filter output are actually edges vs. noise
      - Threshold, Thin



# Canny edge detector

- Filter image with derivative of Gaussian
- Find magnitude and orientation of gradient
- **Non-maximum suppression:**
  - Thin multi-pixel wide “ridges” down to single pixel width
- Linking and thresholding (**hysteresis**):
  - Define two thresholds: low and high
  - Use the high threshold to start edge curves and the low threshold to continue them
- MATLAB: `edge(image, 'canny');`
- `>>help edge`



# The Canny edge detector



**original image (Lena)**

# The Canny edge detector




**magnitude of the gradient**

# The Canny edge detector



thresholding


# The Canny edge detector



Problem:  
pixels along  
this edge  
didn't  
survive the  
thresholding



thinning  
(non-maximum suppression)

# The Canny edge detector



How to turn  
these thick  
regions of  
the gradient  
into  
curves?

# Non-maximum suppression



- Check if pixel is local maximum along gradient direction
  - can require checking interpolated pixels p and r

# The Canny edge detector



**thinning**

(non-maximum suppression)



# Effect of $\sigma$ (Gaussian kernel spread/size)



original



Canny with  $\sigma = 1$



Canny with  $\sigma = 2$

The choice of  $\sigma$  depends on desired behavior

- large  $\sigma$  detects large scale edges
- small  $\sigma$  detects fine features

# So, what scale to choose?


It depends what we're looking for.



Too fine of a scale...can't see the forest for the trees.


Too coarse of a scale...can't tell the maple grain from the cherry.

# Single 2D edge detection filter



Gaussian


$$h_{\sigma}(u, v) = \frac{1}{2\pi\sigma^2} e^{-\frac{u^2+v^2}{2\sigma^2}}$$



derivative of Gaussian

$$\frac{\partial}{\partial x} h_{\sigma}(u, v)$$

Laplacian of Gaussian



$$\nabla^2 h_{\sigma}(u, v)$$

$\nabla^2$  is the **Laplacian** operator:

$$\nabla^2 f = \frac{\partial^2 f}{\partial x^2} + \frac{\partial^2 f}{\partial y^2}$$

# Finish on Thurs...