

CS 4495 Computer Vision

Motion Models

Aaron Bobick
School of Interactive
Computing

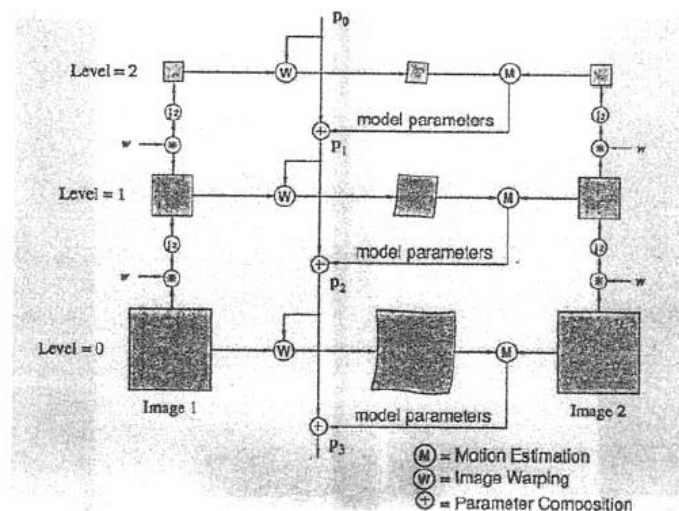


Fig. 1. Diagram of the hierarchical motion estimation framework.

Outline



- Last time: dense motion: optic flow
 - Brightness constancy constraint equation
 - Lucas-Kanade
- 2D Motion models
 - Bergen, '92
 - Pyramids
 - Layers
- Motion fields from 3D motions
- Parametric motion

Visual motion





Many slides adapted from S. Seitz, R. Szeliski, M. Pollefeys, K. Grauman and others...

Motion estimation: Optical flow



Will start by estimating motion of each pixel separately
Then will consider motion of entire image

Problem definition: optical flow


 $I(x, y, t)$

 $I(x, y, t + 1)$

How to estimate pixel motion from image $I(x, y, t)$ to $I(x, y, t + 1)$?

- Solve pixel correspondence problem
 - given a pixel in $I(x, y, t)$, look for nearby pixels of the same color in $I(x, y, t + 1)$

Key assumptions

- **color constancy**: a point in $I(x, y, t)$ looks the same in $I(x, y, t + 1)$
 - For grayscale images, this is brightness constancy
- **small motion**: points do not move very far

This is called the optical flow problem

Optical flow equation

- Combining these two equations

shorthand: $I_x = \frac{\partial I}{\partial x}$

$$\begin{aligned}
 0 &= I(x + u, y + v, t + 1) - I(x, y, t) \\
 &\approx I(x, y, t + 1) + I_x u + I_y v - I(x, y, t) \\
 &\approx [I(x, y, t + 1) - I(x, y, t)] + I_x u + I_y v \\
 &\approx I_t + I_x u + I_y v \\
 &\approx I_t + \nabla I \cdot \langle u, v \rangle
 \end{aligned}$$

In the limit as u and v go to zero, this becomes exact

$$0 = I_t + \nabla I \cdot \langle u, v \rangle$$

Brightness constancy constraint equation

$$I_x u + I_y v + I_t = 0$$

Optical flow equation


$$0 = I_t + \nabla I \cdot \langle u, v \rangle \quad \text{or} \quad I_x u + I_y v + I_t = 0$$

- Q: how many unknowns and equations per pixel?


2 unknowns, one equation

Intuitively, what does this constraint mean?

- The component of the flow in the gradient direction is determined
- The component of the flow parallel to an edge is unknown



Aperture problem



Solving the aperture problem

- How to get more equations for a pixel?
 - Basic idea: impose additional constraints
 - most common is to assume that the flow field is smooth locally
 - one method: pretend the pixel's neighbors have the same (u,v)
 - If we use a 5x5 window, that gives us 25 equations per pixel!

$$0 = I_t(\mathbf{p}_i) + \nabla I(\mathbf{p}_i) \cdot [u \ v]$$

$$\begin{array}{ccc}
 \begin{bmatrix} I_x(\mathbf{p}_1) & I_y(\mathbf{p}_1) \\ I_x(\mathbf{p}_2) & I_y(\mathbf{p}_2) \\ \vdots & \vdots \\ I_x(\mathbf{p}_{25}) & I_y(\mathbf{p}_{25}) \end{bmatrix} & \begin{bmatrix} u \\ v \end{bmatrix} & = - \begin{bmatrix} I_t(\mathbf{p}_1) \\ I_t(\mathbf{p}_2) \\ \vdots \\ I_t(\mathbf{p}_{25}) \end{bmatrix} \\
 \underset{25 \times 2}{A} & \underset{2 \times 1}{d} & \underset{25 \times 1}{b}
 \end{array}$$

Lukas-Kanade flow

- Prob: we have more equations than unknowns

$$\begin{matrix} A & d = b \\ 25 \times 2 & 2 \times 1 \quad 25 \times 1 \end{matrix} \longrightarrow \text{minimize } \|Ad - b\|^2$$

Solution: solve least squares problem

- minimum least squares solution given by solution (in d) of:

$$\begin{matrix} (A^T A) & d = A^T b \\ 2 \times 2 & 2 \times 1 \quad 2 \times 1 \end{matrix}$$

$$\begin{matrix} \begin{bmatrix} \sum I_x I_x & \sum I_x I_y \\ \sum I_x I_y & \sum I_y I_y \end{bmatrix} & \begin{bmatrix} u \\ v \end{bmatrix} = - & \begin{bmatrix} \sum I_x I_t \\ \sum I_y I_t \end{bmatrix} \\ A^T A & & A^T b \end{matrix}$$

- The summations are over all pixels in the K x K window
- This technique was first proposed by Lukas & Kanade (1981)

Eigenvectors of $A^T A$

$$A^T A = \begin{bmatrix} \sum I_x I_x & \sum I_x I_y \\ \sum I_x I_y & \sum I_y I_y \end{bmatrix} = \sum \begin{bmatrix} I_x \\ I_y \end{bmatrix} [I_x \ I_y] = \sum \nabla I (\nabla I)^T$$

- Recall the Harris corner detector: $M = A^T A$ is the *second moment matrix*
- The eigenvectors and eigenvalues of M relate to edge direction and magnitude
 - The eigenvector associated with the larger eigenvalue points in the direction of fastest intensity change
 - The other eigenvector is orthogonal to it

Violating assumptions in Lucas-Kanade

- The motion is large (larger than a pixel)
 - Not-linear: Iterative refinement
 - Local minima: coarse-to-fine estimation
- A point does not move like its neighbors
 - Motion segmentation
- Brightness constancy does not hold
 - Do exhaustive neighborhood search with normalized correlation - tracking features – maybe SIFT – more later....

Violating assumptions in Lucas-Kanade


- The motion is large (larger than a pixel)
 - Not-linear: Iterative refinement
 - Local minima: coarse-to-fine estimation
- A point does not move like its neighbors
 - Motion segmentation
- Brightness constancy does not hold
 - Do exhaustive neighborhood search with normalized correlation - tracking features – maybe SIFT – more later....

Not tangent: Iterative Refinement

Iterative Lukas-Kanade Algorithm


1. Estimate velocity at each pixel by solving Lucas-Kanade equations
2. Warp I_t towards I_{t+1} using the estimated flow field
 - - *use image warping techniques*
3. Repeat until convergence

Optical Flow: Iterative Estimation




(using d for *displacement* here instead of u)


Optical Flow: Iterative Estimation



Optical Flow: Iterative Estimation



Optical Flow: Iterative Estimation



Revisiting the small motion assumption




- Is this motion small enough?
 - Probably not—it's much larger than one pixel (2nd order terms dominate)
 - How might we solve this problem?

Optical Flow: Aliasing




Temporal aliasing causes ambiguities in optical flow because images can have many pixels with the same intensity.

I.e., how do we know which ‘correspondence’ is correct?




To overcome aliasing: coarse-to-fine estimation.

Reduce the resolution!



Coarse-to-fine optical flow estimation




Gaussian pyramid of image 1

$u=1.25$ pixels

$u=2.5$ pixels

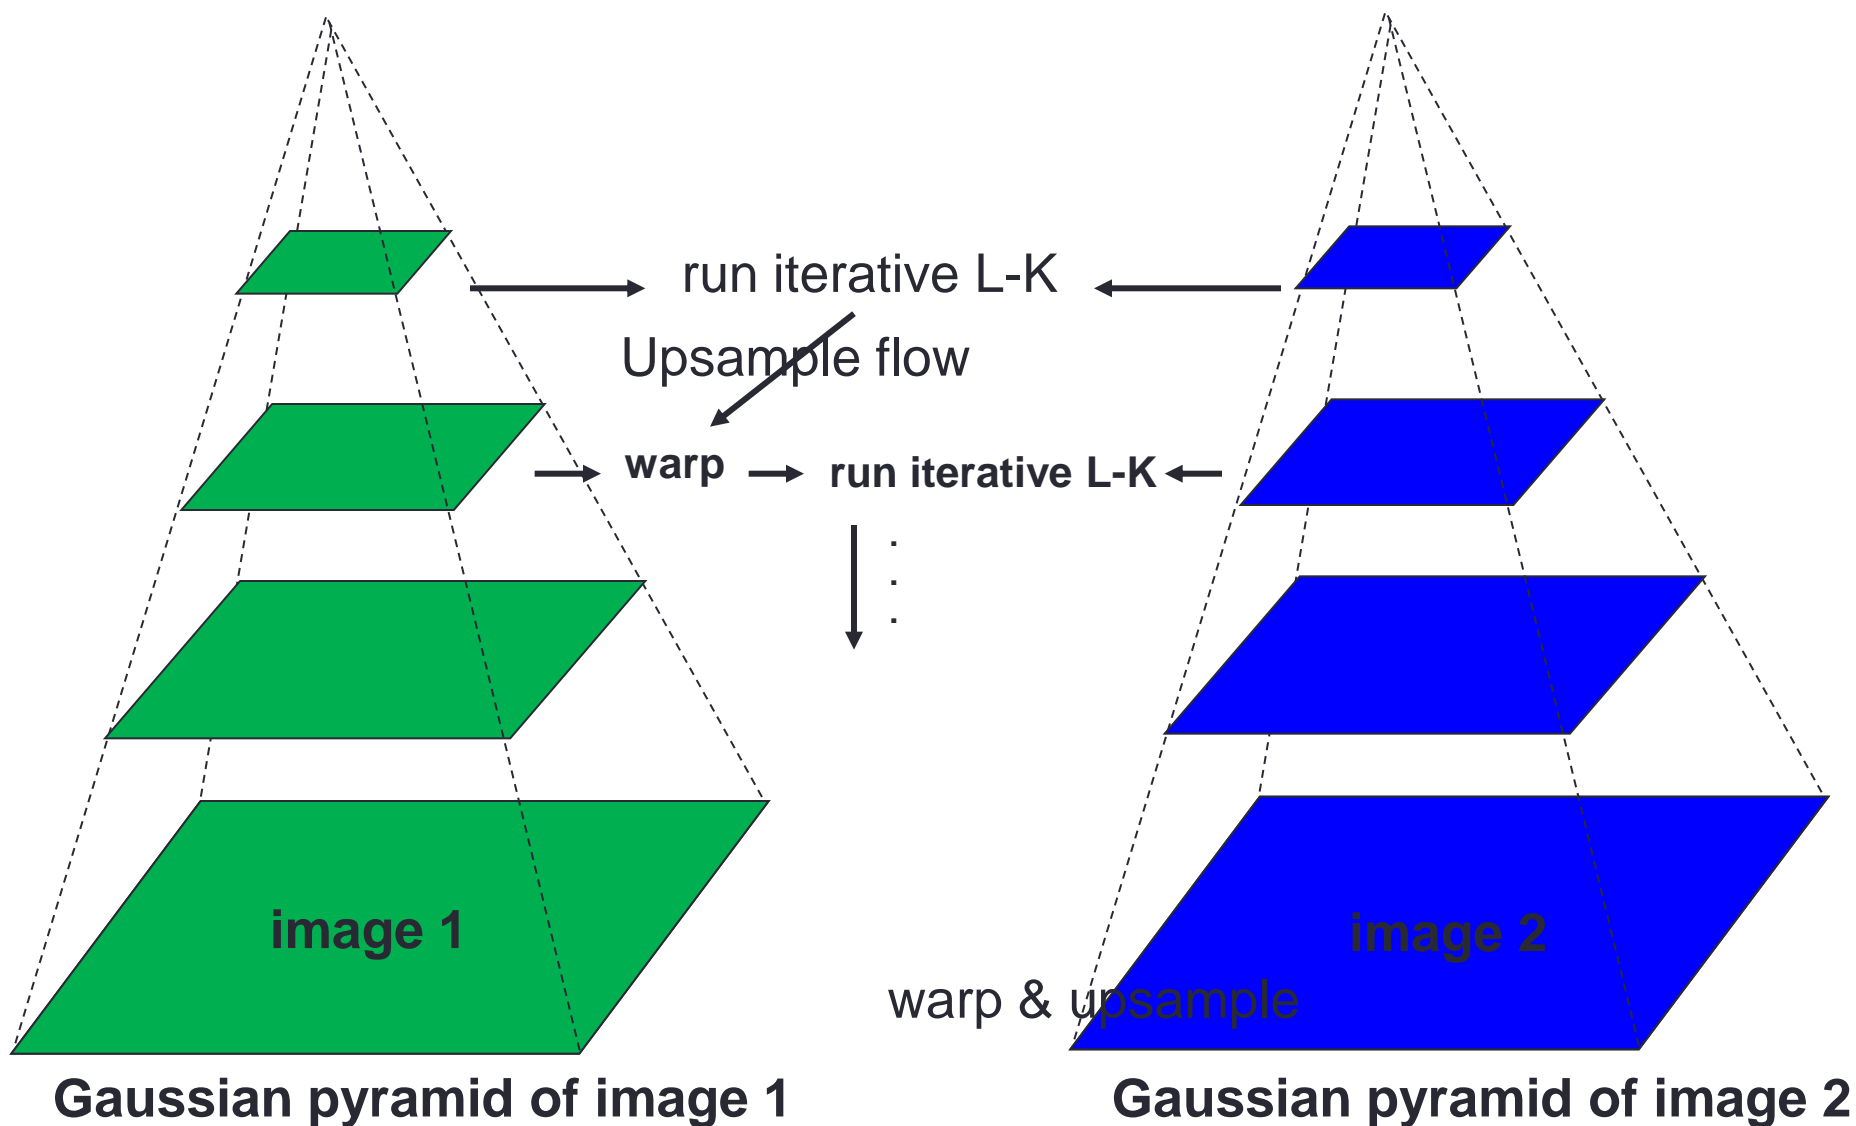
$u=5$ pixels

$u=10$ pixels




Gaussian pyramid of image 2

Coarse-to-fine optical flow estimation



Multi-scale

Remember: Image sub-sampling




1/4



1/8

Throw away every other row and column to create a $1/2$ size image
- called *image sub-sampling*

Bad image sub-sampling



1/2




1/4 (2x zoom)



1/8 (4x zoom)

Aliasing! What do we do?

Gaussian (lowpass) pre-filtering



Gaussian 1/2




G 1/4



G 1/8

Solution: filter the image, *then* subsample

Subsampling with Gaussian pre-filtering



Gaussian $1/2$



G $1/4$



G $1/8$

Band-pass filtering

Gaussian Pyramid (low-pass images)




Laplacian Pyramid (subband images)

These are “bandpass” images (almost).

Band-pass filtering

Gaussian Pyramid (low-pass images)




Laplacian Pyramid (subband images)

These are “bandpass” images (almost).


Laplacian Pyramid

Original
image



- How can we reconstruct (collapse) this pyramid into the original image?


Laplacian Pyramid



- How can we reconstruct (collapse) this pyramid into the original image?

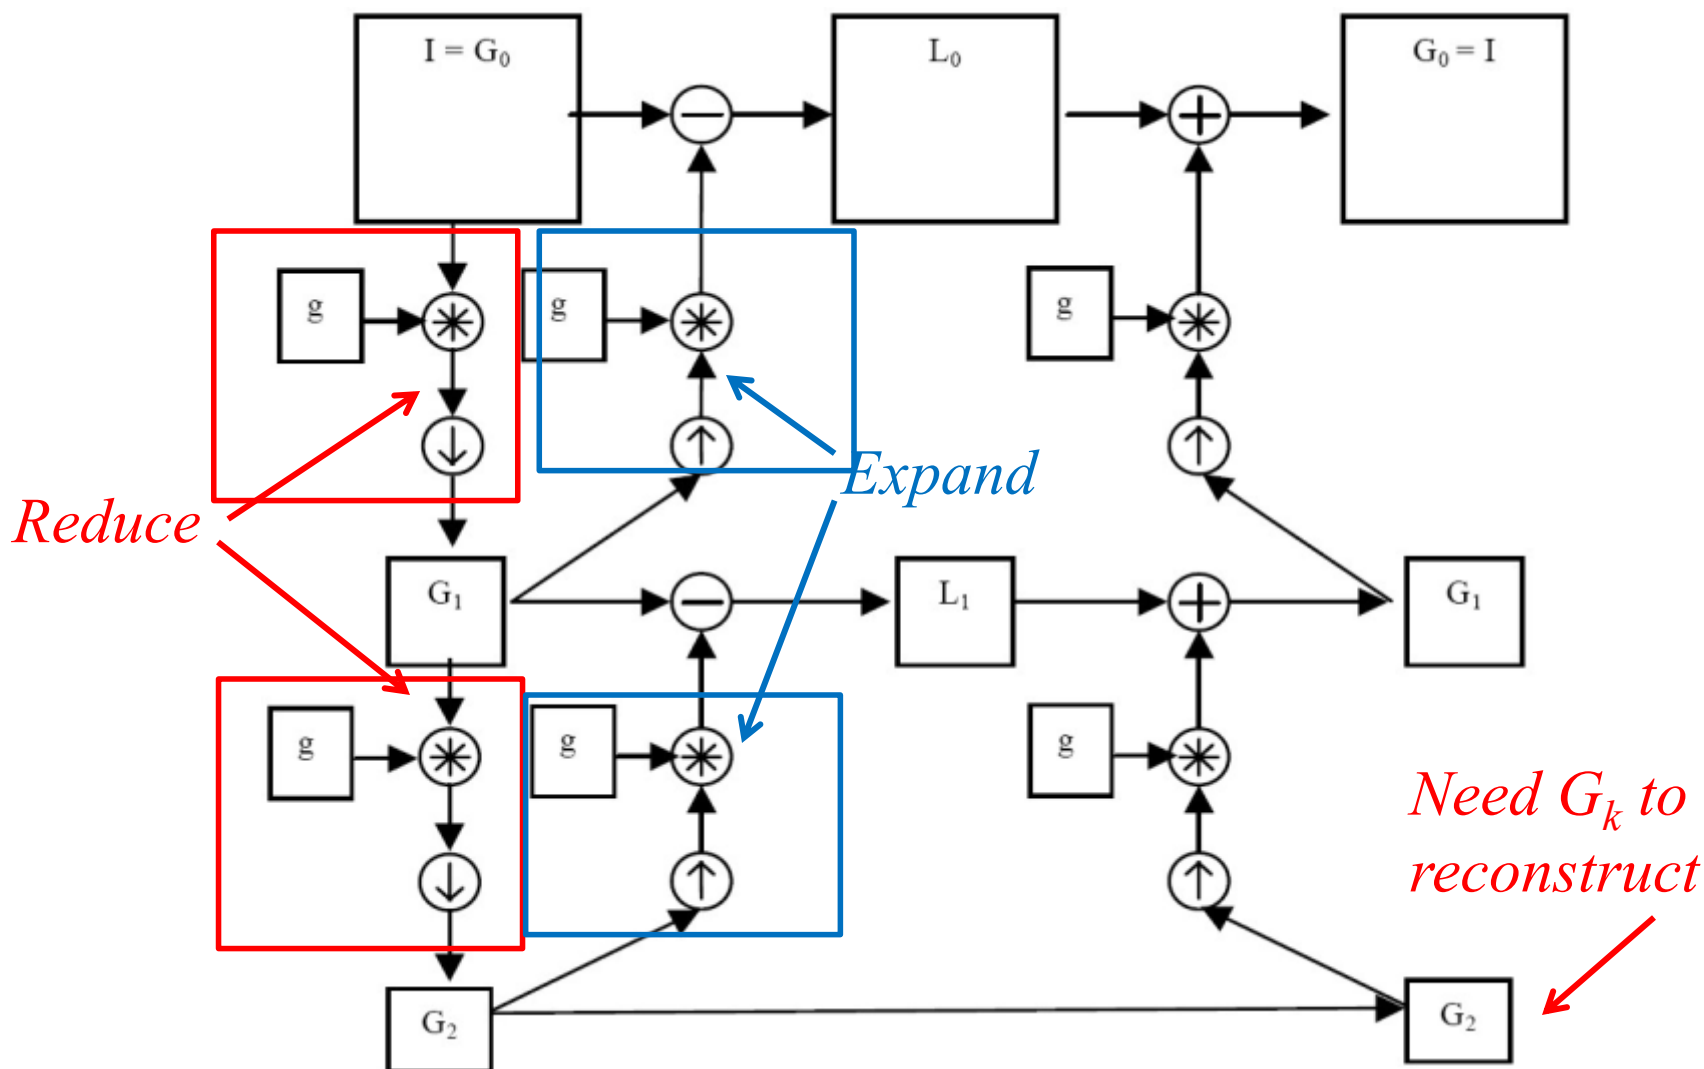
Image Pyramids

Idea: Represent $N \times N$ image as a “pyramid” of $1 \times 1, 2 \times 2, 4 \times 4, \dots, 2^k \times 2^k$ images (assuming $N = 2^k$)




Known as a **Gaussian Pyramid** [Burt and Adelson, 1983]

Computing the Laplacian Pyramid




Reduce and Expand



Reduce


Apply “5-tap” *separable*
filter to make reduced
image.

Reduce and Expand



Reduce

Apply “5-tap” separable filter to make reduced image.




Expand

Apply different “3-tap” separable filters for even and odd pixels to make expanded image...

Just Expand

Apply different “3-tap” separable filters for even and odd pixels to make expanded image.



What can you do with band limited imaged?



Apples and Oranges in bandpass

 L_0 


(a)




(b)




(c)

 L_2 

(d)



(e)



(f)

 L_4 

(g)



(h)



(i)

Reconstructed



(j)



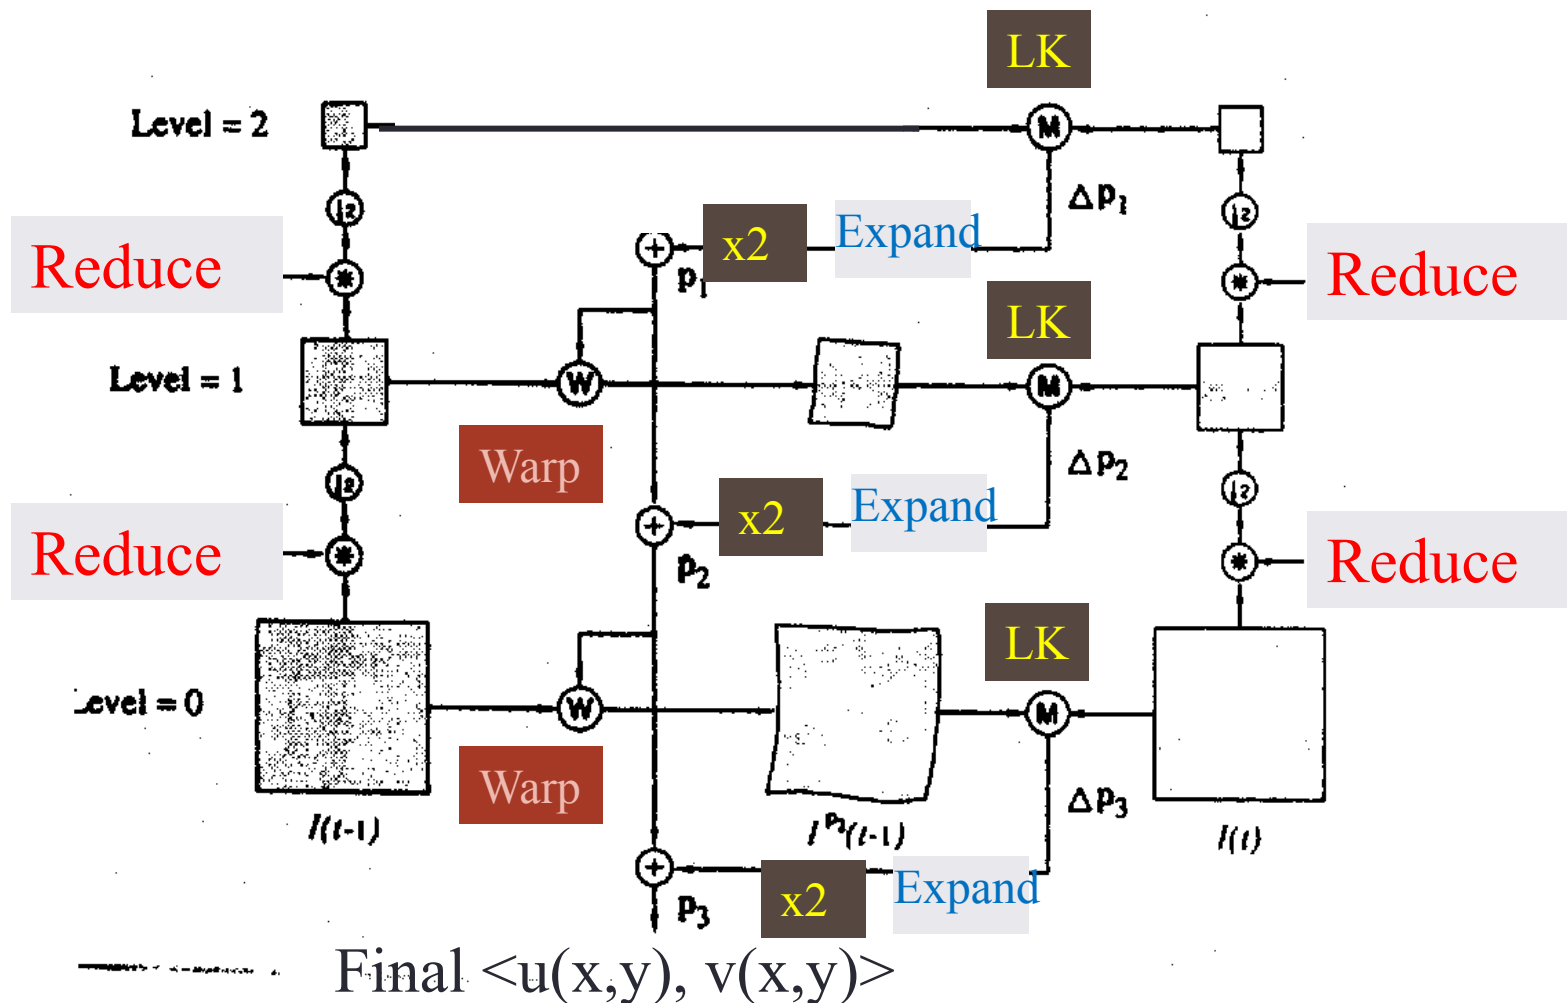
(k)



(l)

Applying pyramids to LK

Coarse-to-fine global motion estimation



Multi-resolution Lucas Kanade Algorithm

Compute Iterative LK at highest level


Initialize $u_{K+1}, v_{K+1} = 0$ at size of level $K+1$

For Each Level i from K to 0


- Upsample u_{i+1}, v_{i+1} to create u_i^p, v_i^p flow fields of now twice resolution as level $i+1$.
- Multiply u_i^p, v_i^p by 2 to get predicted flow
- Warp I_2 according to predicted flow
- Compute I_t –temporal derivative
- Apply LK to get u_i^δ, v_i^δ (the correction in flow)
- Add corrections to obtain the flow $u(i), v(i)$ at i^{th} level, i.e.,

$$u_i = u_i^p + u_i^\delta; v_i = v_i^p + v_i^\delta$$


Optical Flow Results



Optical Flow Results








Lucas-Kanade with Pyramids



Moving to models

- Previous method(s) give dense flow with little or no constraint between locations (smoothness is either explicit or implicit).
- Suppose you “know” that motion is constrained, e.g.
 - Small rotation about horizontal or vertical axis (or both) that is very close to a translation.
 - Distant independent moving objects
- In this case you might “model” the flow...
Ready for another old slide?

Motion models

**Translation****2 unknowns****Similarity****4 unknowns****Affine****6 unknowns****Perspective****8 unknowns**

Focus of Expansion (FOE) - Example

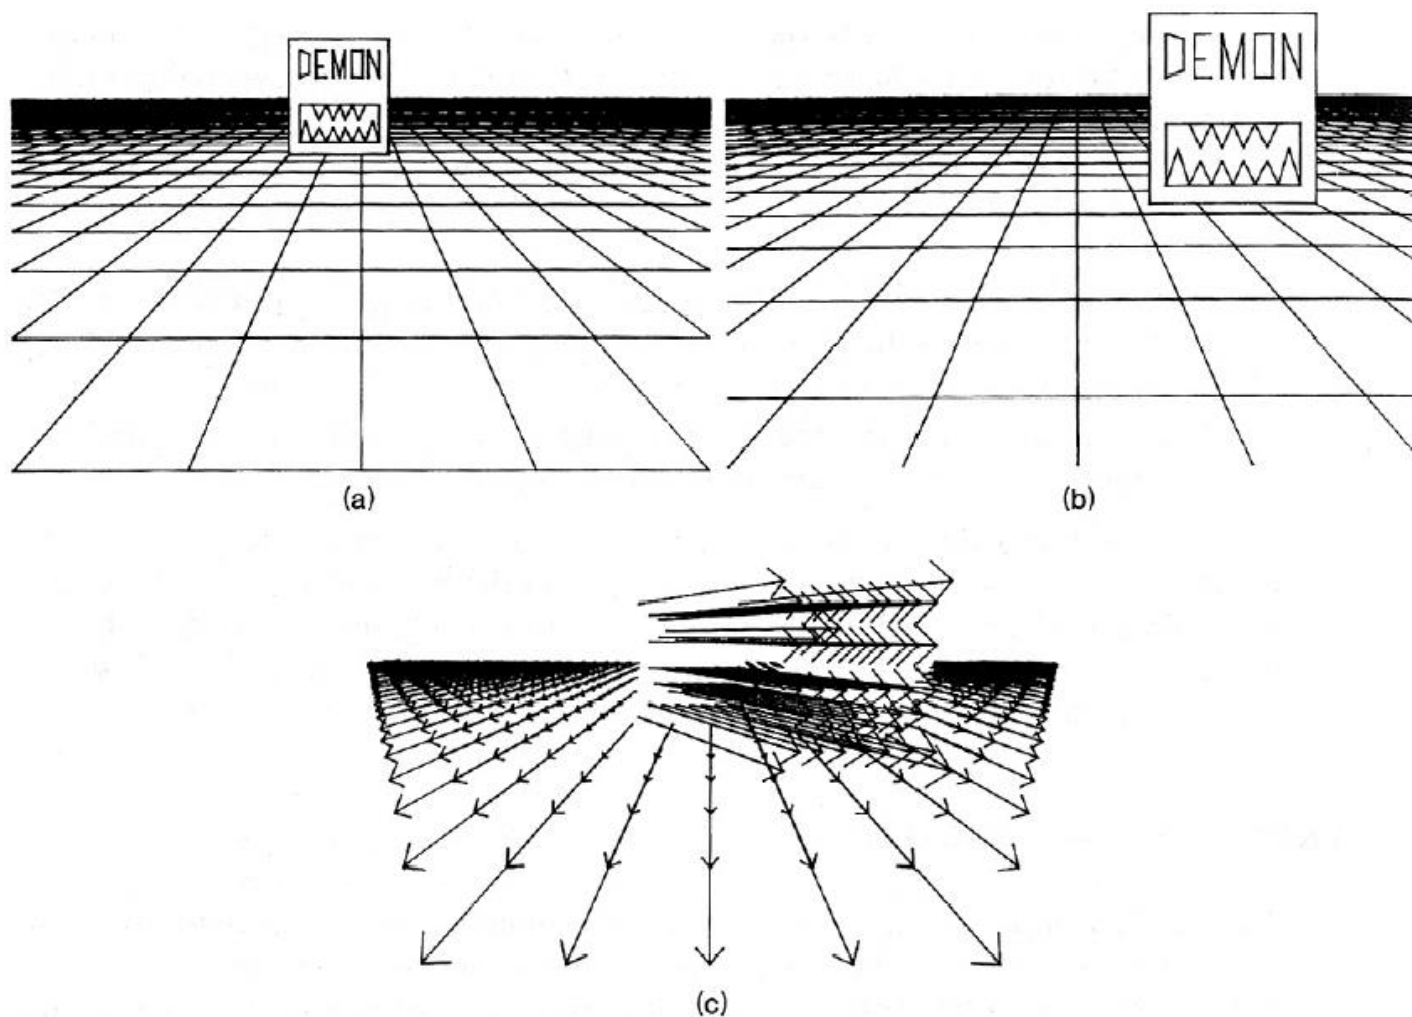


Fig. 7.3 FOE for rectilinear observer motion. (a) An image. (b) Later image. (c) Flow shows different FOEs for static floor and moving object.

Full motion model

- From physics or elsewhere:

$$V = \Omega \times R + T$$

$$\begin{bmatrix} V_x \\ V_y \\ V_z \end{bmatrix} \approx \begin{bmatrix} 0 & -\omega_z & \omega_y \\ \omega_z & 0 & -\omega_x \\ -\omega_y & \omega_x & 0 \end{bmatrix} \begin{bmatrix} X \\ Y \\ Z \end{bmatrix} + \begin{bmatrix} V_{T_x} \\ V_{T_y} \\ V_{T_z} \end{bmatrix}$$

$$\begin{bmatrix} a_x \\ a_y \\ a_z \end{bmatrix} = \begin{bmatrix} 0 & -a_3 & a_2 \\ a_3 & 0 & -a_1 \\ -a_2 & a_1 & 0 \end{bmatrix}$$

$$\begin{bmatrix} V_x \\ V_y \\ V_z \end{bmatrix} = \text{Velocity Vector}$$

$$\begin{bmatrix} V_{T_x} \\ V_{T_y} \\ V_{T_z} \end{bmatrix} = \text{Translational Component of Velocity}$$

$$\begin{bmatrix} \omega_x \\ \omega_y \\ \omega_z \end{bmatrix} = \text{Angular Velocity}$$

General motion

Take derivatives:

$$x = f \frac{X}{Z}$$

$$y = f \frac{Y}{Z}$$

$$u = v_x = f \frac{ZV_x - XV_z}{Z^2} = f \frac{V_x}{Z} - \left(f \frac{X}{Z} \right) \frac{V_z}{Z} = f \frac{V_x}{Z} - x \frac{V_z}{Z}$$

$$v = v_y = f \frac{ZV_y - YV_z}{Z^2} = f \frac{V_y}{Z} - \left(f \frac{Y}{Z} \right) \frac{V_z}{Z} = f \frac{V_y}{Z} - y \frac{V_z}{Z}$$

$$\begin{bmatrix} u(x, y) \\ v(x, y) \end{bmatrix} = \frac{1}{Z(x, y)} \mathbf{A}(x, y) \mathbf{T} + \mathbf{B}(x, y) \mathbf{\Omega}$$

Why is Z only here?

$$\mathbf{A}(x, y) = \begin{bmatrix} -f & 0 & x \\ 0 & -f & y \end{bmatrix} \quad \mathbf{B}(x, y) = \begin{bmatrix} (xy)/f & -(f + x^2)/f & y \\ (f + y^2)/f & -(xy)/f & -x \end{bmatrix}$$

Where \mathbf{T} is translation vector, $\mathbf{\Omega}$ is rotation

If a plane and perspective...

$$aX + bY + cZ + d = 0$$

$$u(x, y) = a_1 + a_2x + a_3y + a_7x^2 + a_8xy$$

$$v(x, y) = a_4 + a_5x + a_6y + a_7xy + a_8y^2$$

If a plane and orthographic...

$$u(x, y) = a_1 + a_2 x + a_3 y$$

$$v(x, y) = a_4 + a_5 x + a_6 y$$

Affine!

Affine motion

$$u(x, y) = a_1 + a_2 x + a_3 y$$

$$v(x, y) = a_4 + a_5 x + a_6 y$$


- Substituting into the brightness constancy equation:

$$I_x \cdot u + I_y \cdot v + I_t \approx 0$$

$$I_x (a_1 + a_2 x + a_3 y) + I_y (a_4 + a_5 x + a_6 y) + I_t \approx 0$$

- Each pixel provides 1 linear constraint in 6 unknowns
- Least squares minimization:

$$Err(\vec{a}) = \sum \left[I_x (a_1 + a_2 x + a_3 y) + I_y (a_4 + a_5 x + a_6 y) + I_t \right]^2$$



Affine motion

- Can sum gradients over window or entire image:

$$Err(\vec{a}) = \sum \left[I_x(a_1 + a_2x + a_3y) + I_y(a_4 + a_5x + a_6y) + I_t \right]^2$$

- Minimize squared error (robustly)

$$\begin{bmatrix} I_x & I_x x_1 & I_x y_1 & I_y & I_y x_1 & I_y y_1 \\ I_x & I_x x_2 & I_x y_2 & I_y & I_y x_2 & I_y y_2 \\ & & \vdots & & & \\ & & \vdots & & & \\ I_x & I_x x_n & I_x y_n & I_y & I_y x_n & I_y y_n \end{bmatrix} \cdot \begin{bmatrix} a_1 \\ a_2 \\ a_3 \\ a_4 \\ a_5 \\ a_6 \end{bmatrix} = - \begin{bmatrix} I_t^1 \\ I_t^2 \\ \vdots \\ I_t^n \end{bmatrix}$$

- This is an example of parametric flow – can substitute any linear model easily. Others with some work.

Hierarchical model-based flow

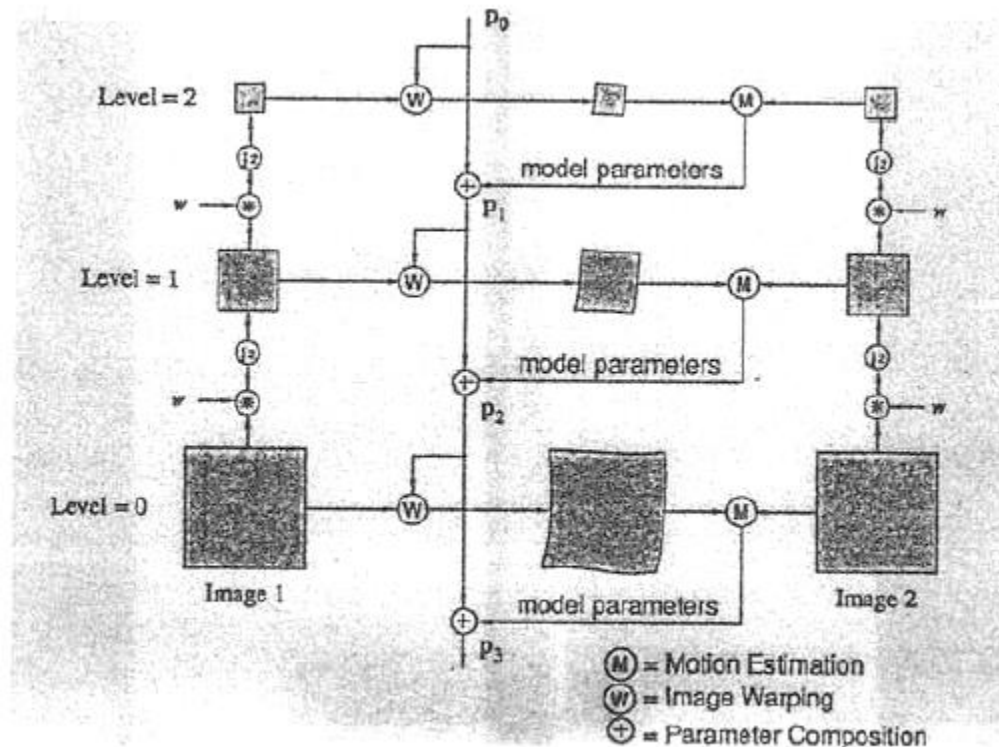



Fig. 1. Diagram of the hierarchical motion estimation framework.

James R. Bergen, P. Anandan, Keith J. Hanna, Rajesh Hingorani:
 "Hierarchical Model-Based Motion Estimation," ECCV 1992: 237-252

Now, if different motion regions...






Layered motion

- Basic idea: break image sequence into “layers” each of which has a coherent motion



What are layers?

- Each layer is defined by an alpha mask and an affine motion model



Motion segmentation with an affine model

$$u(x, y) = a_1 + a_2 x + a_3 y$$

$$v(x, y) = a_4 + a_5 x + a_6 y$$

Local flow
estimates

Motion segmentation with an affine model

$$u(x, y) = a_1 + a_2 x + a_3 y$$

$$v(x, y) = a_4 + a_5 x + a_6 y$$

Equation of a plane
(parameters a_1, a_2, a_3 can be
found by least squares)


Motion segmentation with an affine model

$$u(x, y) = a_1 + a_2 x + a_3 y$$


$$v(x, y) = a_4 + a_5 x + a_6 y$$

Equation of a plane
(parameters a_1, a_2, a_3 can be found by least squares)


1D example




True flow



Local flow estimate



Segmented estimate



Line fitting

“Foreground”






“Background”

Occlusion

How do we estimate the layers?

- Compute local flow in a coarse-to-fine fashion
- Obtain a set of initial affine motion hypotheses
 - Divide the image into blocks and estimate affine motion parameters in each block by least squares
 - Eliminate hypotheses with high residual error
 - Perform k-means clustering on affine motion parameters
 - Merge clusters that are close and retain the largest clusters to obtain a smaller set of hypotheses to describe all the motions in the scene
- Iterate until convergence:
 - Assign each pixel to best hypothesis
 - Pixels with high residual error remain unassigned
 - Perform region filtering to enforce spatial constraints
 - Re-estimate affine motions in each region

Example result



Recovering image motion

- Feature-based methods (e.g. SIFT, Ransac, regression)
 - Extract visual features (corners, textured areas) and track them - sometimes over multiple frames
 - Sparse motion fields, but possibly robust tracking
 - Good for global motion
 - Suitable especially when image motion is large (10-s of pixels)
 - **PS4!**
- Direct-methods (e.g. optical flow)
 - Directly recover image motion from spatio-temporal image brightness variations
 - Dense, local motion fields, but more sensitive to appearance variations
 - Suitable for video and when image motion is small (< 10 pixels)
 - **PS5!!!**

End CS4495