


# CS 4495 Computer Vision

## **RAN**dom **SA**mples **C**onsensus

---

Aaron Bobick  
School of Interactive  
Computing




# Administrivia

- PS 3:
  - Check Piazza - good conversations
  - The F matrix: the actual numbers may vary quite a bit. But check the epipolar lines you get.
    - Normalization: read extra credit part. At least try removing the centroid. Since we're using homogenous coordinates (2D homogenous have 3 elements) it's easy to have a transformation matrix that subtracts off an offset.
- Go back and recheck slides: A 3 vector in these projective geometry is both a point and a line.


# Matching with Features

- Want to compute transformation from one image to the other
- Overall strategy:
  - Compute features
  - Match matching features (duh?)
  - Compute best transformation (translation, affine, homography) from matches



An introductory example:

## ***Harris corner detector***



C.Harris, M.Stephens. “A Combined Corner and Edge Detector”. 1988

# Harris Detector: Mathematics

$$M = A^T A = \begin{bmatrix} \sum I_x I_x & \sum I_x I_y \\ \sum I_x I_y & \sum I_y I_y \end{bmatrix}$$

Measure of corner response:

$$R = \det M - k (\text{trace } M)^2$$


$$\det M = \lambda_1 \lambda_2$$

$$\text{trace } M = \lambda_1 + \lambda_2$$

( $k$  – empirical constant,  $k = 0.04$ - $0.06$ )


# Harris Detector: Mathematics

- $R$  depends only on eigenvalues of  $M$
- $R$  is large for a **corner**
- $R$  is negative with large magnitude for an **edge**
- $|R|$  is small for a **flat** region



# Key point localization

- General idea: find robust extremum (maximum or minimum) both in space and in scale.
- SIFT specific suggestion: use DoG pyramid to find maximum values (remember edge detection?) – then eliminate “edges” and pick only corners.
- More recent: use Harris detector to find maximums in space and then look at the Laplacian pyramid (we’ll do this later) for maximum in scale.




*Each point is compared to its 8 neighbors in the current image and 9 neighbors each in the scales above and below.*



# Point Descriptors


- We know how to detect points
- How to match them? Two parts:
  - Compute a descriptor for each and make the descriptor both as invariant and as distinctive as possible. (Competing goals) SIFT one example.





# Another version of the problem...

Want to find

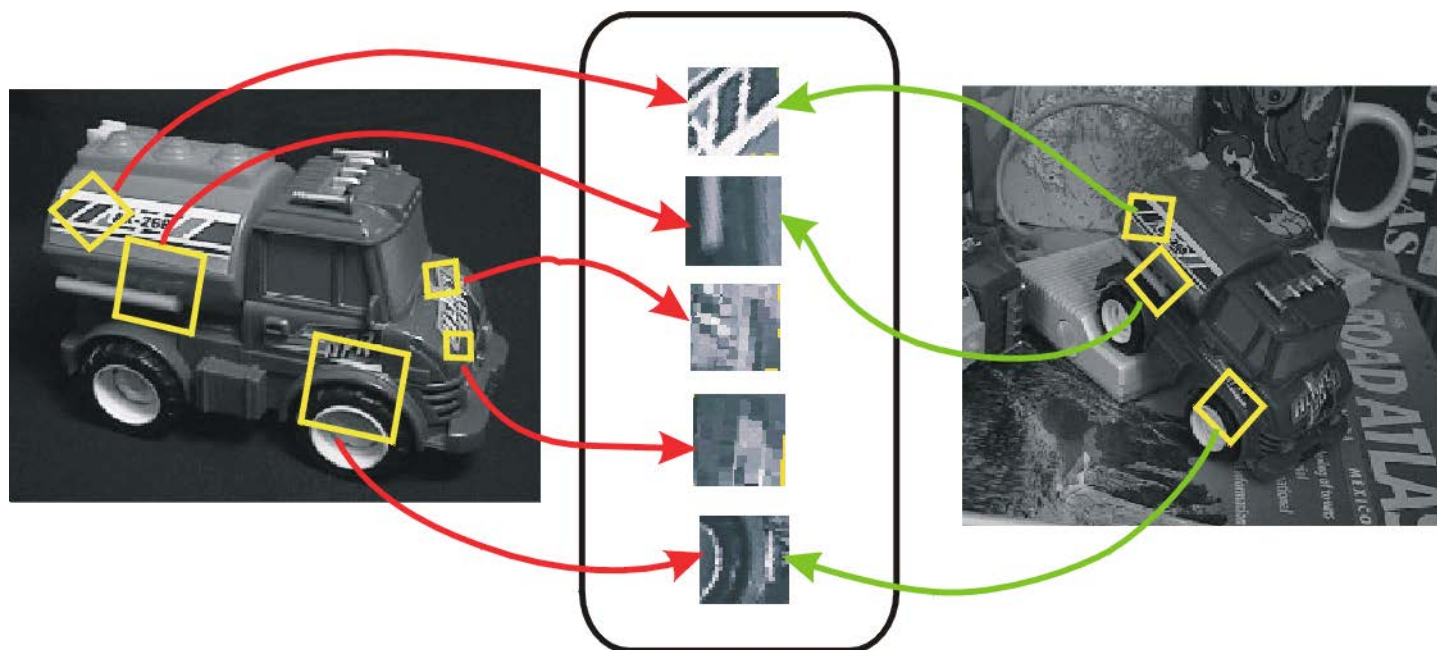


... in here



# Idea of SIFT

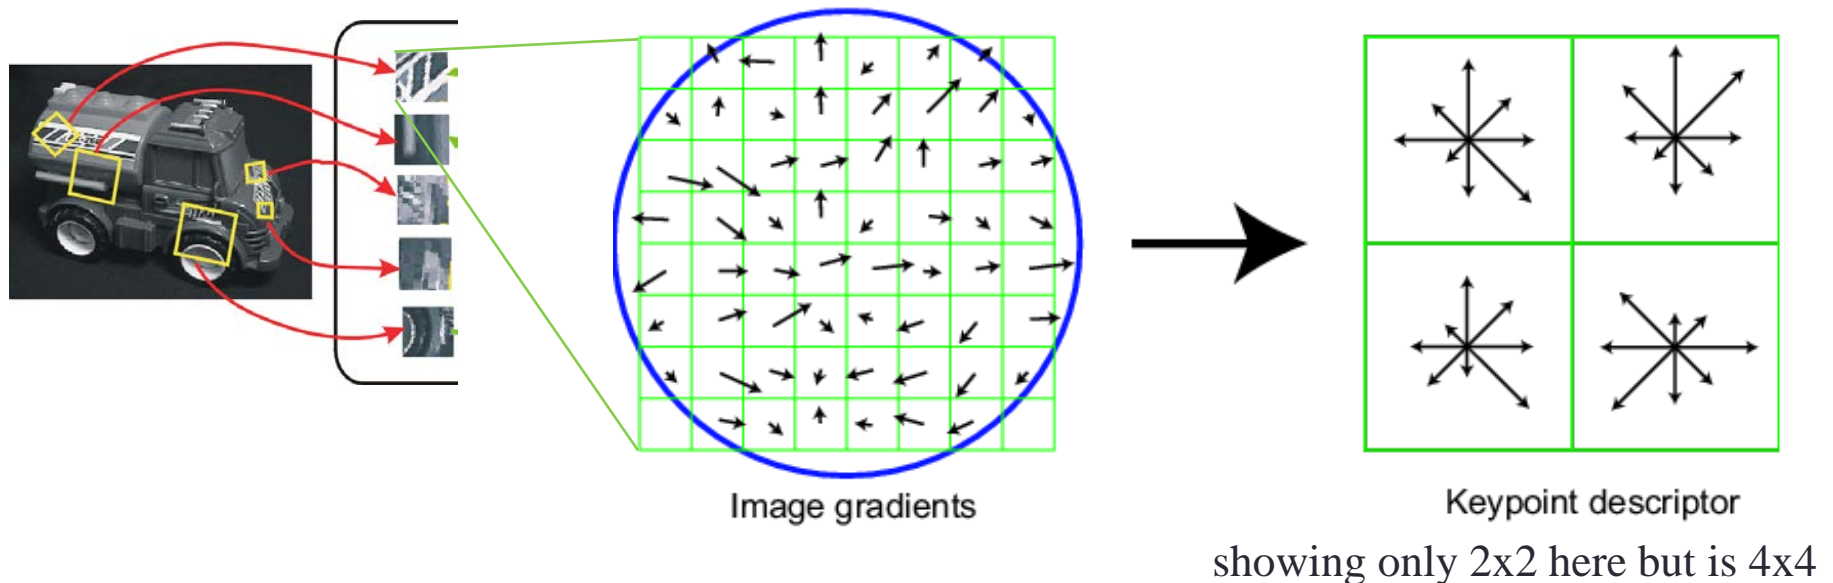
- Image content is transformed into local feature coordinates that are invariant to translation, rotation, scale, and other imaging parameters



**SIFT Features**


# SIFT vector formation

- 4x4 array of gradient orientation histograms over 4x4 pixels
  - not really histogram, weighted by magnitude
- 8 orientations x 4x4 array = 128 dimensions
- Motivation: some sensitivity to spatial layout, but not too much.





# Point Descriptors

- We know how to detect points
- How to match them? Two parts:
  - Compute a descriptor for each and make the descriptor both as invariant and as distinctive as possible. (Competing goals) SIFT one example
  - ***Need to figure out which point matches which..***





# Feature-based alignment outline




- Extract features

# Feature-based alignment outline



- Extract features
- Compute **putative** matches – e.g. “closest descriptor”


# Feature-based alignment outline



- Extract features
- Compute **putative** matches – e.g. “closest descriptor”
- Loop:
  - Hypothesize transformation  $T$  from some matches



# Feature-based alignment outline



- Extract features
- Compute **putative** matches
- Loop:
  - *Hypothesize* transformation  $T$  from some matches
  - *Verify* transformation (search for other matches consistent with  $T$ )

# Feature-based alignment outline



- Extract features
- Compute *putative matches*
- Loop:
  - *Hypothesize* transformation  $T$  from some matches
  - *Verify* transformation (search for other matches consistent with  $T$ )
- Apply transformation


# How to get “putative” matches?

# Feature matching

- Exhaustive search
  - for each feature in one image, look at *all* the other features in the other image(s) – pick best one
- Hashing
  - compute a short descriptor from each feature vector, or hash longer descriptors (randomly)
- Nearest neighbor techniques
  - $k$ -trees and their variants


# Feature-space outlier rejection

- Let's not match all features, but only these that have “similar enough” matches?
- How can we do it?
  - $\text{SSD}(\text{patch1}, \text{patch2}) < \text{threshold}$
  - How to set threshold?



# Feature-space outlier rejection

- A better way [Lowe, 1999]:
  - 1-NN: SSD of the closest match
  - 2-NN: SSD of the second-closest match
  - Look at how much better 1-NN is than 2-NN, e.g. 1-NN/2-NN
  - That is, is our best match so much better than the rest?



# Feature matching

- Exhaustive search
  - for each feature in one image, look at *all* the other features in the other image(s) – pick best one
- Hashing
  - compute a short descriptor from each feature vector, or hash longer descriptors (randomly)
- Nearest neighbor techniques
  - $k$ -trees and their variants
- But...
- Remember: distinctive vs invariant competition? Means:
- *Problem: Even when pick best match, still lots (and lots) of wrong matches – “outliers”*




# Another way to remove mistakes

- Why are we doing matching?
  - To compute a model of the relation between entities
- So this is really “model fitting”

# Fitting

- Choose a parametric model to represent a set of features – *remember this???*



simple model: lines



simple model: circles



complicated model: car

# Fitting: Issues

## Case study: Line detection




- **Noise** in the measured feature locations
- **Extraneous data:** clutter (outliers), multiple lines
- **Missing data:** occlusions

# Typical least squares line fitting

$$E = \sum_{i=1}^n (y_i - mx_i - b)^2$$

- Data:  $(x_1, y_1), \dots, (x_n, y_n)$
- Line equation:  $y_i = mx_i + b$
- Find  $(m, b)$  to minimize



$$E = \sum_{i=1}^n \left( y_i - \begin{bmatrix} x_i & 1 \end{bmatrix} \begin{bmatrix} m \\ b \end{bmatrix} \right)^2 = \left\| \begin{bmatrix} y_1 \\ \vdots \\ y_n \end{bmatrix} - \begin{bmatrix} x_1 & 1 \\ \vdots & \vdots \\ x_n & 1 \end{bmatrix} \begin{bmatrix} m \\ b \end{bmatrix} \right\|^2 = \|\mathbf{y} - \mathbf{X}\mathbf{b}\|^2$$

$$= (\mathbf{y} - \mathbf{X}\mathbf{b})^T (\mathbf{y} - \mathbf{X}\mathbf{b}) = \mathbf{y}^T \mathbf{y} - 2(\mathbf{X}\mathbf{b})^T \mathbf{y} + (\mathbf{X}\mathbf{b})^T (\mathbf{X}\mathbf{b})$$

$$\frac{dE}{d\mathbf{b}} = 2\mathbf{X}^T \mathbf{X}\mathbf{b} - 2\mathbf{X}^T \mathbf{y} = 0$$

$$\mathbf{X}^T \mathbf{X}\mathbf{b} = \mathbf{X}^T \mathbf{y} \Rightarrow \mathbf{b} = (\mathbf{X}^T \mathbf{X})^{-1} \mathbf{X}^T \mathbf{y}$$


*Standard least squares solution to except weird switch of typical names from  $\mathbf{Ax}=\mathbf{b}$*

# Problem with “vertical” least squares

- Not rotation-invariant
- Fails completely for vertical lines


# Total least squares

- Distance between point  $(x_i, y_i)$  and line  $ax+by=d$  ( $a^2+b^2=1$ ):  $|ax_i + by_i - d|$



# Total least squares

- Distance between point  $(x_i, y_i)$  and line  $ax+by=d$
- Find  $(a, b, d)$  to minimize the sum of squared perpendicular distances



$$E = \sum_{i=1}^n (ax_i + by_i - d)^2$$



# Total least squares

- Distance between point  $(x_i, y_i)$  and line  $ax+by=d$
- Find  $(a, b, d)$  to minimize the sum of squared perpendicular distances


$$E = \sum_{i=1}^n (ax_i + by_i - d)^2$$

$$\frac{\partial E}{\partial d} = \sum_{i=1}^n -2(ax_i + by_i - d) = 0$$

$$E = \sum_{i=1}^n (a(x_i - \bar{x}) + b(y_i - \bar{y}))^2 = \left\| \begin{bmatrix} x_1 - \bar{x} & y_1 - \bar{y} \\ \vdots & \vdots \\ x_n - \bar{x} & y_n - \bar{y} \end{bmatrix} \begin{bmatrix} a \\ b \end{bmatrix} \right\|^2 = (UN)^T (UN)$$

$$\frac{dE}{dN} = 2(U^T U)N = 0$$

Solution to  $(U^T U)N = 0$ , subject to  $\|N\|^2 = 1$ : eigenvector of  $U^T U$  associated with the smallest eigenvalue (Again SVD to least squares solution to *homogeneous linear system*  $UN = 0$ )



$$d = \frac{a}{n} \sum_{i=1}^n x_i + \frac{b}{n} \sum_{i=1}^n y_i = a\bar{x} + b\bar{y}$$

# Least squares as likelihood maximization


- **Generative model:** line points are corrupted by Gaussian noise in the direction perpendicular to the line

$$\begin{pmatrix} x \\ y \end{pmatrix} = \begin{pmatrix} u \\ v \end{pmatrix} + \varepsilon \begin{pmatrix} a \\ b \end{pmatrix}$$

point  
on the  
line


noise:  
sampled from  
zero-mean  
Gaussian with  
std. dev.  $\sigma$

normal  
direction




# Least squares: Robustness to (very) non-Gaussian noise

- Least squares fit to the red points:



# Least squares: Robustness to (very) non-Gaussian noise

- Least squares fit with an outlier:




Problem: squared error heavily penalizes outliers

# Robust estimators

- General approach: minimize 
$$\sum_i \rho(r_i(x_i, \theta); \sigma)$$


$r_i(x_i, \theta)$  – residual of  $i$ th point w.r.t. model parameters  $\theta$

$\rho$  – robust function with scale parameter  $\sigma$




The robust function  $\rho$  behaves like squared distance for small values of the residual  $u$  but saturates for larger values of  $u$

# Choosing the scale: Just right



The effect of the outlier is minimized


# Choosing the scale: Too small



The error value is almost the same for every point and the fit is very poor



# Choosing the scale: Too large




Behaves much the same as least squares

# “Find consistent matches”???

- Some points (many points) are static in the world
- Some are not
- Need to find the right ones so can compute pose.
- Well tried approach:
  - Random Sample Consensus (RANSAC)

# Simpler Example

- Fitting a straight line



# Discard Outliers


- No point with  $d > t$
- RANSAC:
  - RANdom SAmple Consensus
  - Fischler & Bolles 1981
  - Copes with a large proportion of outliers

M. A. Fischler, R. C. Bolles. [Random Sample Consensus: A Paradigm for Model Fitting with Applications to Image Analysis and Automated Cartography](#). Comm. of the ACM, Vol 24, pp 381-395, 1981.

# RANSAC

(**RAN**dom **SA**mples **C**onsensus) :

Fischler & Bolles in '81.




## Algorithm:

1. **Sample** (randomly) the number of points required to fit the model
2. **Solve** for model parameters using sample
3. **Score** by the fraction of *inliers* within a preset threshold of the model

**Repeat** 1-3 until the best model is found with high confidence

# RANSAC

Line fitting example




Algorithm:

1. **Sample** (randomly) the number of points required to fit the model ( $\#=2$ )
2. **Solve** for model parameters using sample
3. **Score** by the fraction of inliers within a preset threshold of the model

**Repeat** 1-3 until the best model is found with high confidence

# RANSAC

Line fitting example



Algorithm:


1. **Sample** (randomly) the number of points required to fit the model ( $\#=2$ )
2. **Solve** for model parameters using sample
3. **Score** by the fraction of inliers within a preset threshold of the model

**Repeat** 1-3 until the best model is found with high confidence

# RANSAC

Line fitting example

$$N_I = 6$$



Algorithm:


1. **Sample** (randomly) the number of points required to fit the model ( $\#=2$ )
2. **Solve** for model parameters using the sample
3. **Score** by the fraction of inliers within a preset threshold of the model

**Repeat** 1-3 until the best model is found with high confidence



# RANSAC

Line fitting example



Algorithm:

1. **Sample** (randomly) the number of points required to fit the model ( $\#=2$ )
2. **Solve** for model parameters using sample
3. **Score** by the fraction of inliers within a preset threshold of the model

**Repeat** 1-3 until the best model is found with high confidence

# Best Line has most support

- More support -> better fit

# RANSAC for general model

- A given model has a *minimal set* – the smallest number of samples from which the model can be computed.
  - Line: 2 points
- Image transformations are models. Minimal set of  $s$  of point pairs/matches:
  - Translation: pick one point pair
  - Homography (for plane) – pick 4 point pairs
  - Fundamental matrix – pick 8 point pairs (really 7 but lets not go there)
- Algorithm
  - Randomly select  $s$  points (or point pairs) to form a sample
  - Instantiate a model
  - Get consensus set  $S_i$
  - If  $|S_i| > T$ , terminate and return model
  - Repeat for  $N$  trials, return model with  $\max |S_i|$

# Distance Threshold

- Requires noise distribution
- If **Location**: Gaussian noise with  $\sigma^2 = 1$
- Then **Distance**  $d$  has **Chi** distribution with  $k$  degrees of freedoms
- If one dimension, e.g. distance off a line, then 1DOF

$$f(t) = \frac{\sqrt{2}e^{-\frac{t^2}{2}}}{\sqrt{\pi}}, t \geq 0$$

- For 95% cumulative threshold when Gaussian with  $\sigma^2$ :

$$t^2 = 3.84\sigma^2$$

- That is: if  $t^2 = 3.84\sigma^2$  then 95% prob that  $d < t$  when point is inlier

# How many samples ?

- We want: at least one sample with all inliers
- Can't guarantee: probability  $p$  e.g.  $p = 0.99$

# Choosing the parameters

- Initial number of points  $s$ 
  - Typically minimum number needed to fit the model
- Distance threshold  $t$ 
  - Choose  $t$  so probability for inlier is high (e.g. 0.95)
  - Zero-mean Gaussian noise with std. dev.  $\sigma$ :  $t^2 = 3.84\sigma^2$
- Number of samples  $N$ 
  - Choose  $N$  so that, with probability  $p$ , at least one random sample is free from outliers (e.g.  $p = 0.99$ ) (outlier ratio:  $e$ )

# Calculate N

- $s$  – number of points to compute solution
- $p$  – probability of success
- $e$  – proportion outliers, so % inliers =  $(1 - e)$
- $P(\text{sample set with all inliers}) = (1 - e)^s$
- $P(\text{sample set will have at least one outlier}) = (1 - (1 - e)^s)$
- $P(\text{all } N \text{ samples have outlier}) = (1 - (1 - e)^s)^N$
- We want  $P(\text{all } N \text{ samples have outlier}) < (1 - p)$
- So  $(1 - (1 - e)^s)^N < 1 - p$

$$N > \log(1 - p) / \log(1 - (1 - e)^s)$$



# Samples required for inliers only in a sample

- Set  $p=0.99$  – chance of getting good sample


$s = 2, e = 5\%$	$\Rightarrow N=2$
$s = 2, e = 50\%$	$\Rightarrow N=17$
$s = 4, e = 5\%$	$\Rightarrow N=3$
$s = 4, e = 50\%$	$\Rightarrow N=72$
$s = 8, e = 5\%$	$\Rightarrow N=5$
$s = 8, e = 50\%$	$\Rightarrow N=1177$

s	proportion of outliers $e$						
	5%	10%	20%	25%	30%	40%	50%
2	2	3	5	6	7	11	17
3	3	4	7	9	11	19	35
4	3	5	9	13	17	34	72
5	4	6	12	17	26	57	146
6	4	7	16	24	37	97	293
7	4	8	20	33	54	163	588
8	5	9	26	44	78	272	1177


- $N$  increases steeply with  $s$

$$N > \log(1-p) / \log(1-(1-e)^s)$$


$$N = f(e), \text{ *not the number of points!* }$$



# Matching features




# Random Sample Consensus (1)







## Random Sample Consensus (2)



# Least squares fit




# RANSAC for estimating homography


- RANSAC loop:
    1. Select four feature pairs (at random)
    2. Compute homography  $\mathbf{H}$  (exact)
    3. Compute *inliers* where  $SSD(p_i', \mathbf{H} p_i) < \varepsilon$
    4. Keep largest set of inliers
    5. Re-compute least-squares  $\mathbf{H}$  estimate on all of the inliers
- 

# 2D transformation models


- Similarity  
(translation, scale, rotation)



- Affine



- Projective  
(homography)





# Adaptively determining the number of samples

- Inlier ratio  $e$  is often unknown a priori, so pick worst case, e.g. 50%, and adapt if more inliers are found, e.g. 80% would yield  $e=0.2$
- Adaptive procedure:
  - $N = \infty$ , sample\_count = 0,  $e = 1.0$
  - While  $N > \text{sample\_count}$ 
    - Choose a sample and count the number of inliers
    - Set  $e_0 = 1 - (\text{number of inliers})/(\text{total number of points})$
    - If  $e_0 < e$  Set  $e = e_0$  and recompute  $N$  from  $e$ :
$$N = \log(1 - p) / \log(1 - (1 - e)^s)$$
  - Increment the sample\_count by 1

# RANSAC conclusions

## Good

- Simple and general
- Applicable to many different problems, often works well in practice
- Robust to outliers
- Applicable for larger number of parameters than Hough transform
- Parameters are easier to choose than Hough transform

## Bad

- Computational time grows quickly number of parameters
- Not as good for getting multiple fits
- Really not good for approximate models

## Common applications

- Computing a homography (e.g., image stitching)
- Estimating fundamental matrix (relating two views)
- Every problem in robot vision