

# CS 4495 Computer Vision

## *Tracking 1- Kalman, Gaussian*

---

Aaron Bobick  
School of Interactive  
Computing



# Administrivia


- PS5 – will be out this Thurs
  - Due **Sun** Nov 10<sup>th</sup> 11:55pm
- Calendar (tentative) done for the year
  - PS6: 11/14, due 11/24   PS7: 11/26, due 12/5
  - EXAM: Tues before Thanksgiving. Covers concepts and basics
  - So no final on Dec 12....

# Tracking

- Slides “adapted” from Kristen Grauman, Deva Ramanan, but mostly from Svetlana Lazebnik

# Some examples

- Older examples:



- State of the art:

<http://www.youtube.com/watch?v=InqV34BcheM>

# Feature tracking

- So far, we have only considered optical flow estimation in a pair of images
- If we have more than two images, we can compute the optical flow from each frame to the next
- Given a point in the first image, we can in principle reconstruct its path by simply “following the arrows”

# Tracking challenges

- Ambiguity of optical flow
  - Find good features to track
- Large motions
  - Discrete search instead of Lucas-Kanade
- Changes in shape, orientation, color
  - Allow some matching flexibility
- Occlusions, disocclusions
  - Need mechanism for deleting, adding new features
- Drift – errors may accumulate over time
  - Need to know when to terminate a track

# Handling large displacements

- Define a small area around a pixel as the template
- Match the template against each pixel within a search area in next image – just like stereo matching!
- Use a match measure such as SSD or correlation
- After finding the best discrete location, can use Lucas-Kanade to get sub-pixel estimate (think of the template as the coarse level of the pyramid).

# Tracking over many frames

- Select features in first frame
- For each frame:
  - Update positions of tracked features
    - Discrete search or Lucas-Kanade
  - Start new tracks if needed
  - Terminate inconsistent tracks
    - Compute similarity with corresponding feature in the previous frame or in the first frame where it's visible
- This is done by many companies and systems – often ad hoc rules tailored to the context.



# Shi-Tomasi feature tracker

- Find good features using eigenvalues of second-moment matrix – *you've seen this now twice!*
  - Key idea: “good” features to track are the ones that can be tracked reliably
- From frame to frame, track with Lucas-Kanade and a pure translation model
  - More robust for small displacements, can be estimated from smaller neighborhoods
- Check consistency of tracks by affine registration to the first observed instance of the feature
  - Affine model is more accurate for larger displacements
  - Comparing to the first frame helps to minimize drift

# Tracking example



Figure 1: Three frame details from Woody Allen's *Manhattan*. The details are from the 1st, 11th, and 21st frames of a subsequence from the movie.




Figure 2: The traffic sign windows from frames 1,6,11,16,21 as tracked (top), and warped by the computed deformation matrices (bottom).

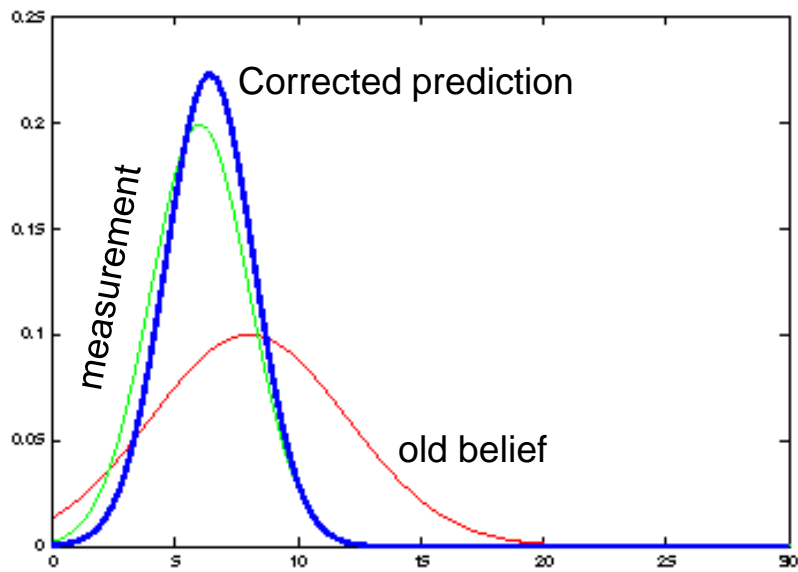

# Tracking with dynamics

- Key idea: Given a model of expected motion, predict where objects will occur in next frame, even before seeing the image
  - Restrict search for the object
  - Improved estimates since measurement noise is reduced by trajectory smoothness

# Tracking as inference

- The hidden state consists of the true parameters we care about, denoted  $X$ .
- The measurement is our noisy observation that results from the underlying state, denoted  $Y$ .
- At each time step, state changes (from  $X_{t-1}$  to  $X_t$ ) and we get a new observation  $Y_t$ .
- Our goal: recover most likely state  $X_t$  given
  - All observations seen so far.
  - Knowledge about dynamics of state transitions.

# Tracking as inference: intuition



Time t

Time t+1

# Steps of tracking

- **Prediction:** What is the next state of the object given past measurements?

$$P(X_t | Y_0 = y_0, \dots, Y_{t-1} = y_{t-1})$$

# Steps of tracking

- **Prediction:** What is the next state of the object given past measurements?

$$P(X_t | Y_0 = y_0, \dots, Y_{t-1} = y_{t-1})$$

- **Correction:** Compute an updated estimate of the state from prediction and measurements

$$P(X_t | Y_0 = y_0, \dots, Y_{t-1} = y_{t-1}, Y_t = y_t)$$

# Steps of tracking

- Prediction: What is the next state of the object given past measurements?

$$P(X_t | Y_0 = y_0, \dots, Y_{t-1} = y_{t-1})$$

- Correction: Compute an updated estimate of the state from prediction and measurements (posterior)

$$P(X_t | Y_0 = y_0, \dots, Y_{t-1} = y_{t-1}, Y_t = y_t)$$

- *Tracking can be seen as the process of propagating the **posterior** distribution of state given measurements across time*



# Simplifying assumptions

- Only the immediate past matters

$$P(X_t | X_0, \dots, X_{t-1}) = \boxed{P(X_t | X_{t-1})}$$

dynamics model

# Simplifying assumptions

- Only the immediate past matters

$$P(X_t | X_0, \dots, X_{t-1}) = \boxed{P(X_t | X_{t-1})}$$

dynamics model

- Measurements depend only on the current state

$$P(Y_t | X_0, Y_0, \dots, X_{t-1}, Y_{t-1}, X_t) = \boxed{P(Y_t | X_t)}$$

observation model

# Simplifying assumptions

- Only the immediate past matters


$$P(X_t | X_0, \dots, X_{t-1}) = P(X_t | X_{t-1})$$

dynamics model

- Measurements depend only on the current state


$$P(Y_t | X_0, Y_0, \dots, X_{t-1}, Y_{t-1}, X_t) = P(Y_t | X_t)$$

observation model



# Tracking as induction

- Base case:
  - Assume we have initial prior that predicts state in absence of any evidence:  $P(X_0)$
  - At the first frame, *correct* this given the value of  $Y_0=y_0$
- Given corrected estimate for frame  $t$ :
  - Predict for frame  $t+1$
  - Correct for frame  $t+1$



# Tracking as induction

- Base case:
  - Assume we have initial prior that predicts state in absence of any evidence:  $P(X_0)$
  - At the first frame, *correct* this given the value of  $Y_0=y_0$

$$P(X_0 | Y_0 = y_0) = \frac{P(y_0 | X_0)P(X_0)}{P(y_0)} \propto P(y_0 | X_0)P(X_0)$$

# Prediction

- Prediction involves guessing  $P(X_t | y_0, \dots, y_{t-1})$   
given  $P(X_{t-1} | y_0, \dots, y_{t-1})$

# Prediction

- Prediction involves guessing  $P(X_t | y_0, \dots, y_{t-1})$   
given  $P(X_{t-1} | y_0, \dots, y_{t-1})$

$$P(X_t | y_0, \dots, y_{t-1}) \\ = \int P(X_t, X_{t-1} | y_0, \dots, y_{t-1}) dX_{t-1}$$

Law of total probability - Marginalization

# Prediction

- Prediction involves guessing  $P(X_t | y_0, \dots, y_{t-1})$   
given  $P(X_{t-1} | y_0, \dots, y_{t-1})$

$$\begin{aligned} P(X_t | y_0, \dots, y_{t-1}) \\ &= \int P(X_t, X_{t-1} | y_0, \dots, y_{t-1}) dX_{t-1} \\ &= \int P(X_t | X_{t-1}, y_0, \dots, y_{t-1}) P(X_{t-1} | y_0, \dots, y_{t-1}) dX_{t-1} \end{aligned}$$

Conditioning on  $X_{t-1}$



# Prediction

- Prediction involves guessing  $P(X_t | y_0, \dots, y_{t-1})$   
given  $P(X_{t-1} | y_0, \dots, y_{t-1})$

$$\begin{aligned} &P(X_t | y_0, \dots, y_{t-1}) \\ &= \int P(X_t, X_{t-1} | y_0, \dots, y_{t-1}) dX_{t-1} \\ &= \int P(X_t | X_{t-1}, y_0, \dots, y_{t-1}) P(X_{t-1} | y_0, \dots, y_{t-1}) dX_{t-1} \\ &= \int P(X_t | X_{t-1}) P(X_{t-1} | y_0, \dots, y_{t-1}) dX_{t-1} \end{aligned}$$

Independence assumption

# Correction

- Correction involves computing  $P(X_t | y_0, \dots, y_t)$  given predicted value  $P(X_t | y_0, \dots, y_{t-1})$  and  $y_t$

# Correction

- Correction involves computing  $P(X_t | y_0, \dots, y_t)$  given predicted value  $P(X_t | y_0, \dots, y_{t-1})$  and  $y_t$

$$\begin{aligned} P(X_t | y_0, \dots, y_t) \\ = \frac{P(y_t | X_t, y_0, \dots, y_{t-1}) P(X_t | y_0, \dots, y_{t-1})}{P(y_t | y_0, \dots, y_{t-1})} \end{aligned}$$

Bayes rule

# Correction

- Correction involves computing  $P(X_t | y_0, \dots, y_t)$  given predicted value  $P(X_t | y_0, \dots, y_{t-1})$  and  $y_t$

$$P(X_t | y_0, \dots, y_t)$$

$$= \frac{P(y_t | X_t, y_0, \dots, y_{t-1}) P(X_t | y_0, \dots, y_{t-1})}{P(y_t | y_0, \dots, y_{t-1})}$$

$$= \frac{P(y_t | X_t) P(X_t | y_0, \dots, y_{t-1})}{P(y_t | y_0, \dots, y_{t-1})}$$

**Independence assumption**  
(observation  $y_t$  depends only on state  $X_t$ )

# Correction

- Correction involves computing  $P(X_t | y_0, \dots, y_t)$  given predicted value  $P(X_t | y_0, \dots, y_{t-1})$  and  $y_t$

$$\begin{aligned}
 &P(X_t | y_0, \dots, y_t) \\
 &= \frac{P(y_t | X_t, y_0, \dots, y_{t-1})P(X_t | y_0, \dots, y_{t-1})}{P(y_t | y_0, \dots, y_{t-1})} \\
 &= \frac{P(y_t | X_t)P(X_t | y_0, \dots, y_{t-1})}{P(y_t | y_0, \dots, y_{t-1})} \\
 &= \frac{P(y_t | X_t)P(X_t | y_0, \dots, y_{t-1})}{\int P(y_t | X_t)P(X_t | y_0, \dots, y_{t-1})dX_t}
 \end{aligned}$$

Really a  
normalization

Conditioning on  $X_t$

# Summary: Prediction and correction

- Prediction:

$$P(X_t | y_0, \dots, y_{t-1}) = \int \underbrace{P(X_t | X_{t-1})}_{\text{dynamics model}} \underbrace{P(X_{t-1} | y_0, \dots, y_{t-1})}_{\text{corrected estimate from previous step}} dX_{t-1}$$

# Summary: Prediction and correction

- Prediction:

$$P(X_t | y_0, \dots, y_{t-1}) = \int \underbrace{P(X_t | X_{t-1})}_{\text{dynamics model}} \underbrace{P(X_{t-1} | y_0, \dots, y_{t-1})}_{\text{corrected estimate from previous step}} dX_{t-1}$$

- Correction:

$$P(X_t | y_0, \dots, y_t) = \frac{\underbrace{P(y_t | X_t)}_{\text{observation model}} \underbrace{P(X_t | y_0, \dots, y_{t-1})}_{\text{predicted estimate}}}{\int P(y_t | X_t) P(X_t | y_0, \dots, y_{t-1}) dX_t}$$

# Linear Dynamic Models

- Dynamics model: state undergoes linear transformation plus Gaussian noise

- $$\mathbf{x}_t \sim N \left( D_t \mathbf{x}_{t-1}, \Sigma_{d_t} \right)$$

- Observation model: measurement is linearly transformed state plus Gaussian noise

$$\mathbf{y}_t \sim N \left( M_t \mathbf{x}_t, \Sigma_{m_t} \right)$$



# Example: Constant velocity (1D)

- State vector is position and velocity

$$x_t = \begin{bmatrix} p_t \\ v_t \end{bmatrix} \quad \begin{aligned} p_t &= p_{t-1} + (\Delta t)v_{t-1} + \varepsilon \\ v_t &= v_{t-1} + \xi \end{aligned} \quad \begin{array}{l} \text{(greek letters} \\ \text{denote noise} \\ \text{terms)} \end{array}$$

$$x_t = D_t x_{t-1} + \text{noise} = \begin{bmatrix} 1 & \Delta t \\ 0 & 1 \end{bmatrix} \begin{bmatrix} p_{t-1} \\ v_{t-1} \end{bmatrix} + \text{noise}$$

- Measurement is position only

$$y_t = Mx_t + \text{noise} = \begin{bmatrix} 1 & 0 \end{bmatrix} \begin{bmatrix} p_t \\ v_t \end{bmatrix} + \text{noise}$$

# Example: Constant acceleration (1D)

- State vector is position, velocity, and acceleration

$$x_t = \begin{bmatrix} p_t \\ v_t \\ a_t \end{bmatrix} \quad \begin{aligned} p_t &= p_{t-1} + (\Delta t)v_{t-1} + \varepsilon \\ v_t &= v_{t-1} + (\Delta t)a_{t-1} + \xi \\ a_t &= a_{t-1} + \zeta \end{aligned} \quad \begin{array}{l} \text{(greek letters} \\ \text{denote noise} \\ \text{terms)} \end{array}$$

$$x_t = D_t x_{t-1} + \text{noise} = \begin{bmatrix} 1 & \Delta t & 0 \\ 0 & 1 & \Delta t \\ 0 & 0 & 1 \end{bmatrix} \begin{bmatrix} p_{t-1} \\ v_{t-1} \\ a_{t-1} \end{bmatrix} + \text{noise}$$


- Measurement is position only

$$y_t = M x_t + \text{noise} = \begin{bmatrix} 1 & 0 & 0 \end{bmatrix} \begin{bmatrix} p_t \\ v_t \\ a_t \end{bmatrix} + \text{noise}$$


# The Kalman filter

- Method for tracking linear dynamical models in Gaussian noise
- The predicted/corrected state distributions are Gaussian
  - You only need to maintain the mean and covariance
  - The calculations are easy (all the integrals can be done in closed form)

# Propagation of Gaussian densities



# The Kalman Filter: 1D state



# 1D Kalman filter: Prediction

- Linear dynamic model defines predicted state evolution, with noise

- $$X_t \sim N(dx_{t-1}, \sigma_d^2)$$

- Want to estimate distribution for next predicted state

$$P(X_t | y_0, \dots, y_{t-1}) = \int P(X_t | X_{t-1}) P(X_{t-1} | y_0, \dots, y_{t-1}) dX_{t-1}$$

# 1D Kalman filter: Prediction

- Linear dynamic model defines predicted state evolution, with noise

$$X_t \sim N(dx_{t-1}, \sigma_d^2)$$

- Want to estimate distribution for next predicted state

$$P(X_t | y_0, \dots, y_{t-1}) = N(\mu_t^-, (\sigma_t^-)^2)$$

- Update the mean:  $\mu_t^- = d\mu_{t-1}^+$
- Update the variance:  $(\sigma_t^-)^2 = \sigma_d^2 + (d\sigma_{t-1}^+)^2$

# 1D Kalman filter: Correction

- Mapping of state to measurements:  $Y_t \sim N(mx_t, \sigma_m^2)$
- Predicted state:  $P(X_t | y_0, \dots, y_{t-1}) = N(\mu_t^-, (\sigma_t^-)^2)$
- Want to estimate corrected distribution

$$P(X_t | y_0, \dots, y_t) = \frac{P(y_t | X_t)P(X_t | y_0, \dots, y_{t-1})}{\int P(y_t | X_t)P(X_t | y_0, \dots, y_{t-1})dX_t}$$



# 1D Kalman filter: Correction

- Mapping of state to measurements:  $Y_t \sim N(mx_t, \sigma_m^2)$
- Predicted state:  $P(X_t | y_0, \dots, y_{t-1}) = N(\mu_t^-, (\sigma_t^-)^2)$
- We define the corrected distribution to be:

$$P(X_t | y_0, \dots, y_t) \equiv N(\mu_t^+, (\sigma_t^+)^2)$$

# 1D Kalman filter: Correction

- Mapping of state to measurements:  $Y_t \sim N(mx_t, \sigma_m^2)$

- Predicted state:  $P(X_t | y_0, \dots, y_{t-1}) = N(\mu_t^-, (\sigma_t^-)^2)$

- Want to estimate corrected distribution

$$P(X_t | y_0, \dots, y_t) = N(\mu_t^+, (\sigma_t^+)^2)$$

- Update the mean: 
$$\mu_t^+ = \frac{\mu_t^- \sigma_m^2 + m y_t (\sigma_t^-)^2}{\sigma_m^2 + m^2 (\sigma_t^-)^2}$$

- Update the variance:

$$(\sigma_t^+)^2 = \frac{\sigma_m^2 (\sigma_t^-)^2}{\sigma_m^2 + m^2 (\sigma_t^-)^2}$$

# 1D Kalman filter: Correction

From:

$$\mu_t^+ = \frac{\mu_t^- \sigma_m^2 + m y_t (\sigma_t^-)^2}{\sigma_m^2 + m^2 (\sigma_t^-)^2}$$

Prediction of  $x$

Measurement guess of  $x$

$$\mu_t^+ = \frac{\mu_t^- \sigma_m^2}{\sigma_m^2 + m^2 (\sigma_t^-)^2} + \frac{y_t}{m} \frac{(\sigma_t^-)^2}{\sigma_m^2 + m^2 (\sigma_t^-)^2}$$

Variance of prediction

Variance of  $x$  computed from the measurement

- What is this?
- ***The weighted average of prediction and measurement based on variances!***

# Prediction vs. correction

$$\mu_t^+ = \frac{\mu_t^- \sigma_m^2 + m y_t (\sigma_t^-)^2}{\sigma_m^2 + m^2 (\sigma_t^-)^2} \quad (\sigma_t^+)^2 = \frac{\sigma_m^2 (\sigma_t^-)^2}{\sigma_m^2 + m^2 (\sigma_t^-)^2}$$

- What if there is no prediction uncertainty  $(\sigma_t^- = 0)$ ?

- $$\mu_t^+ = \mu_t^- \quad (\sigma_t^+)^2 = 0$$

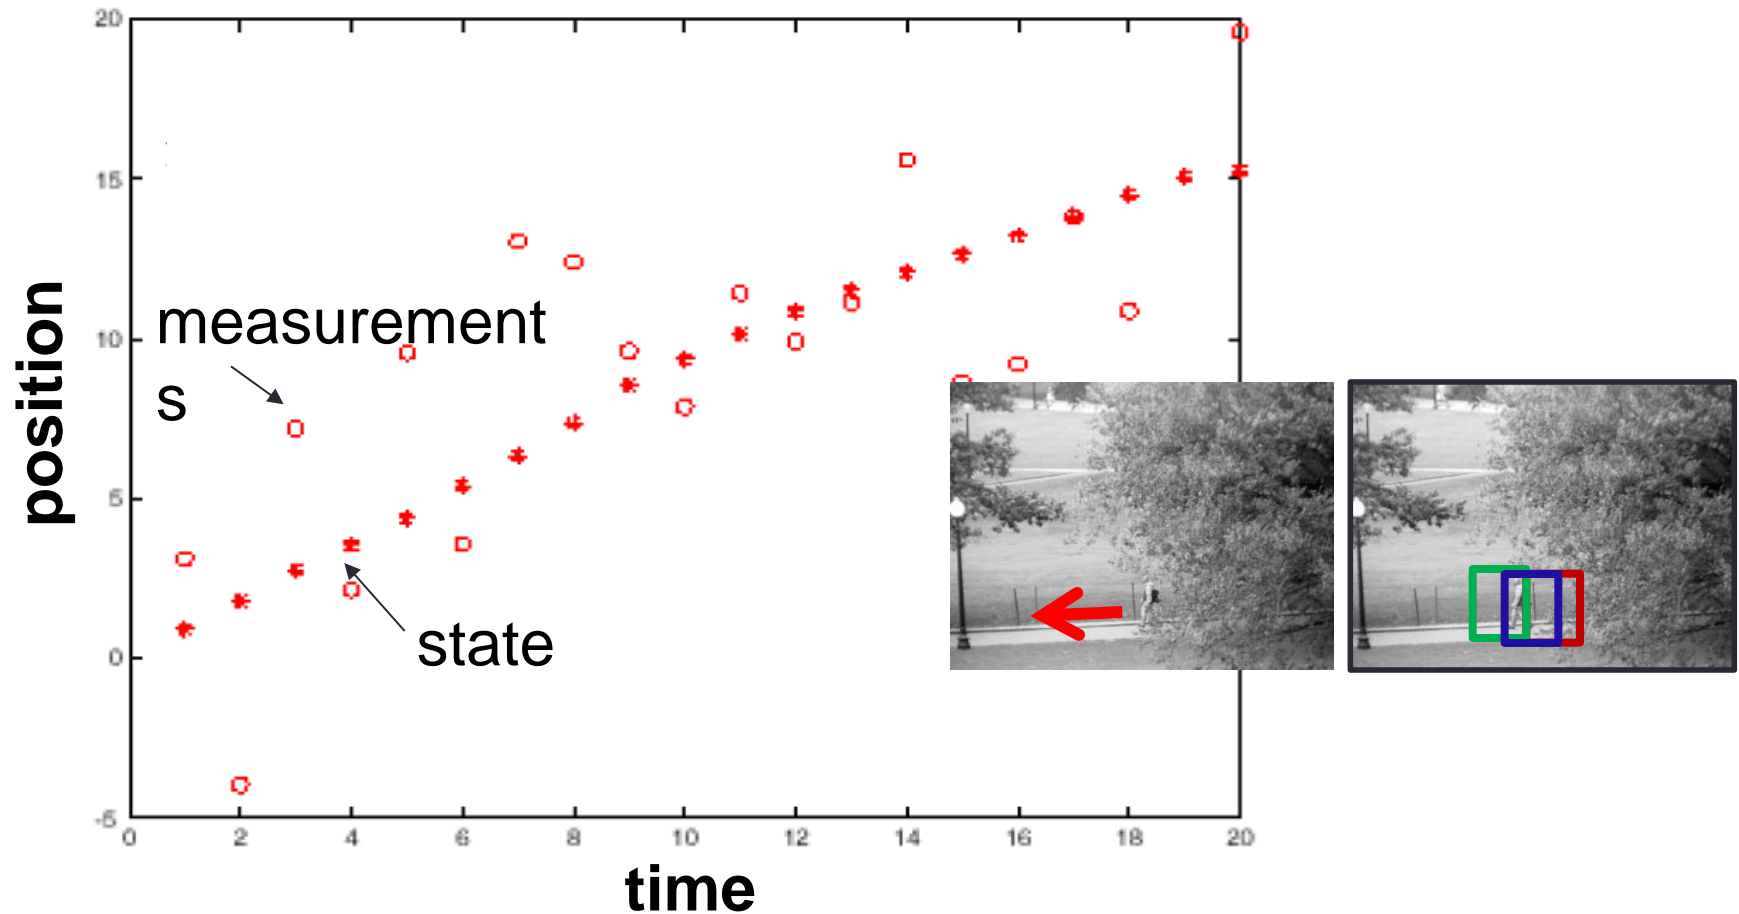
The measurement is ignored!

- What if there is no measurement uncertainty  $(\sigma_m = 0)$ ?

$$\mu_t^+ = \frac{y_t}{m} \quad (\sigma_t^+)^2 = 0$$


The prediction is ignored!

# Recall: constant velocity example

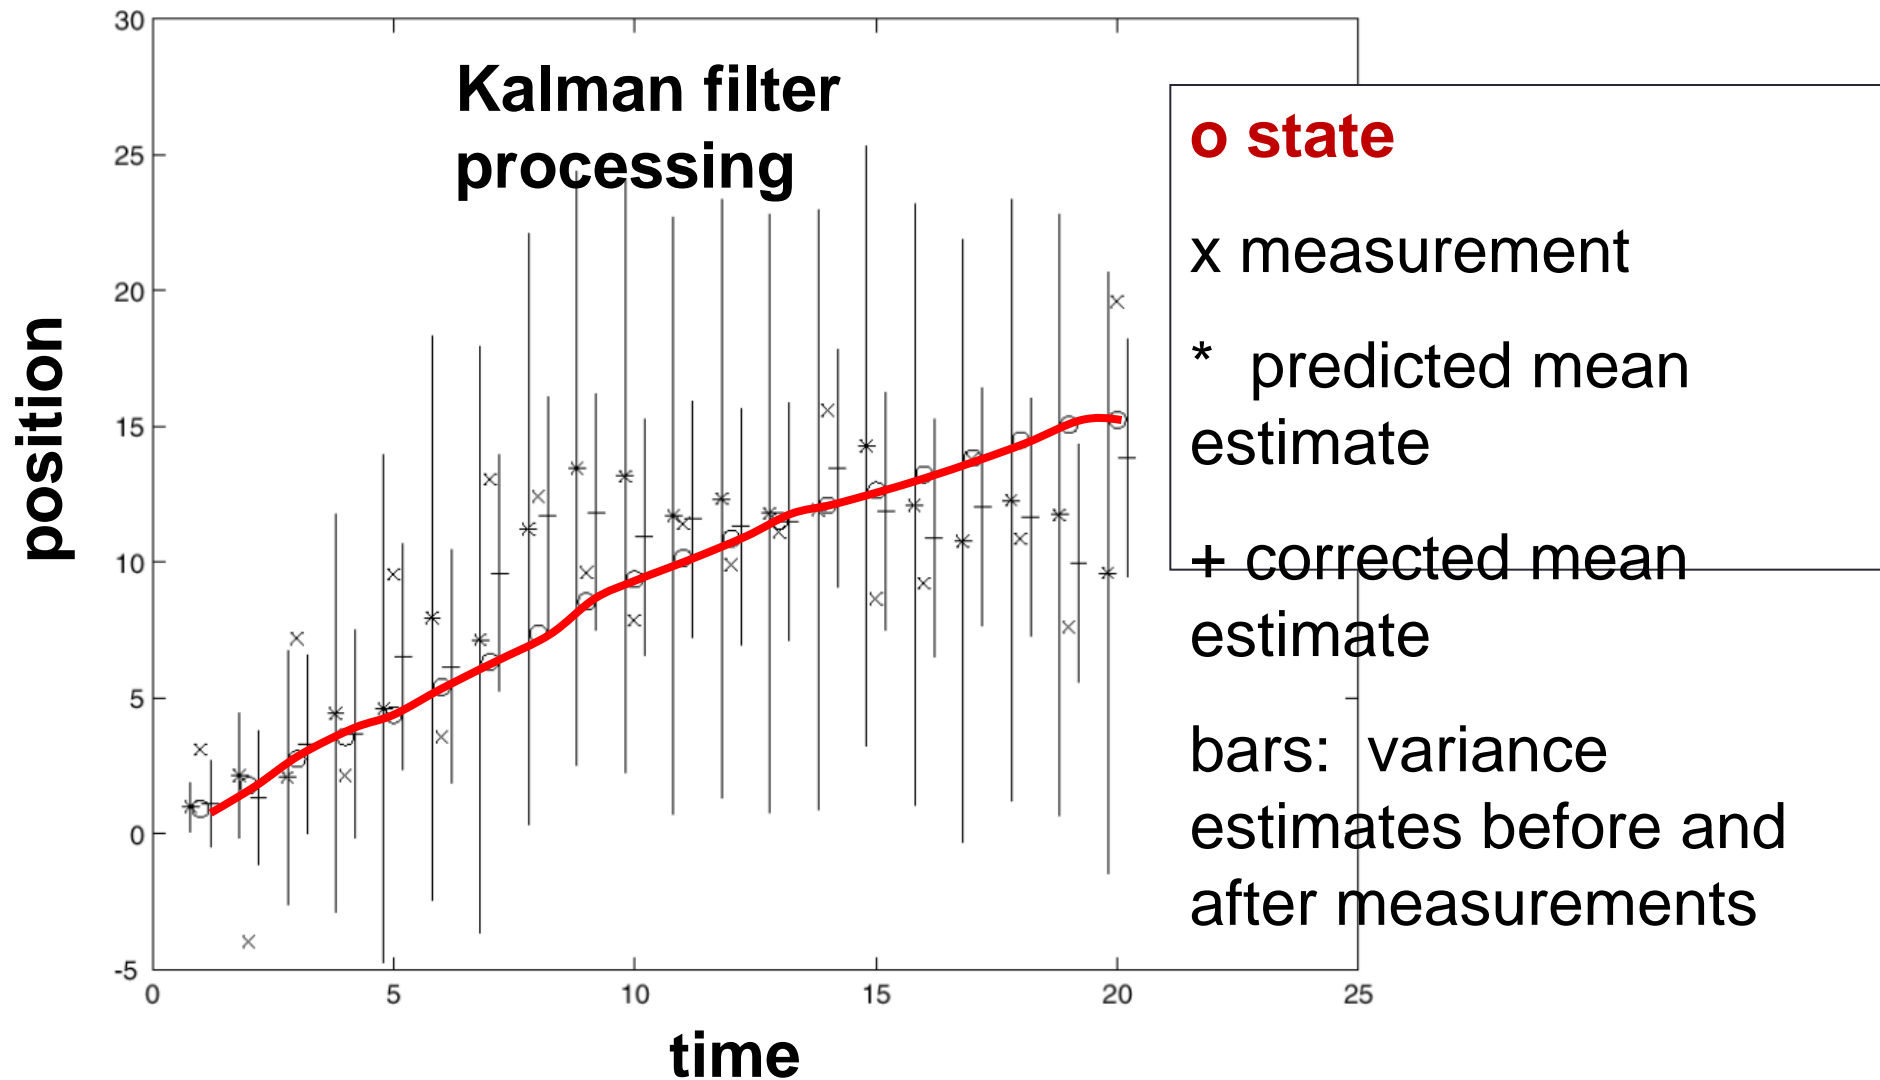


State is 2d: position + velocity  
Measurement is 1d: position

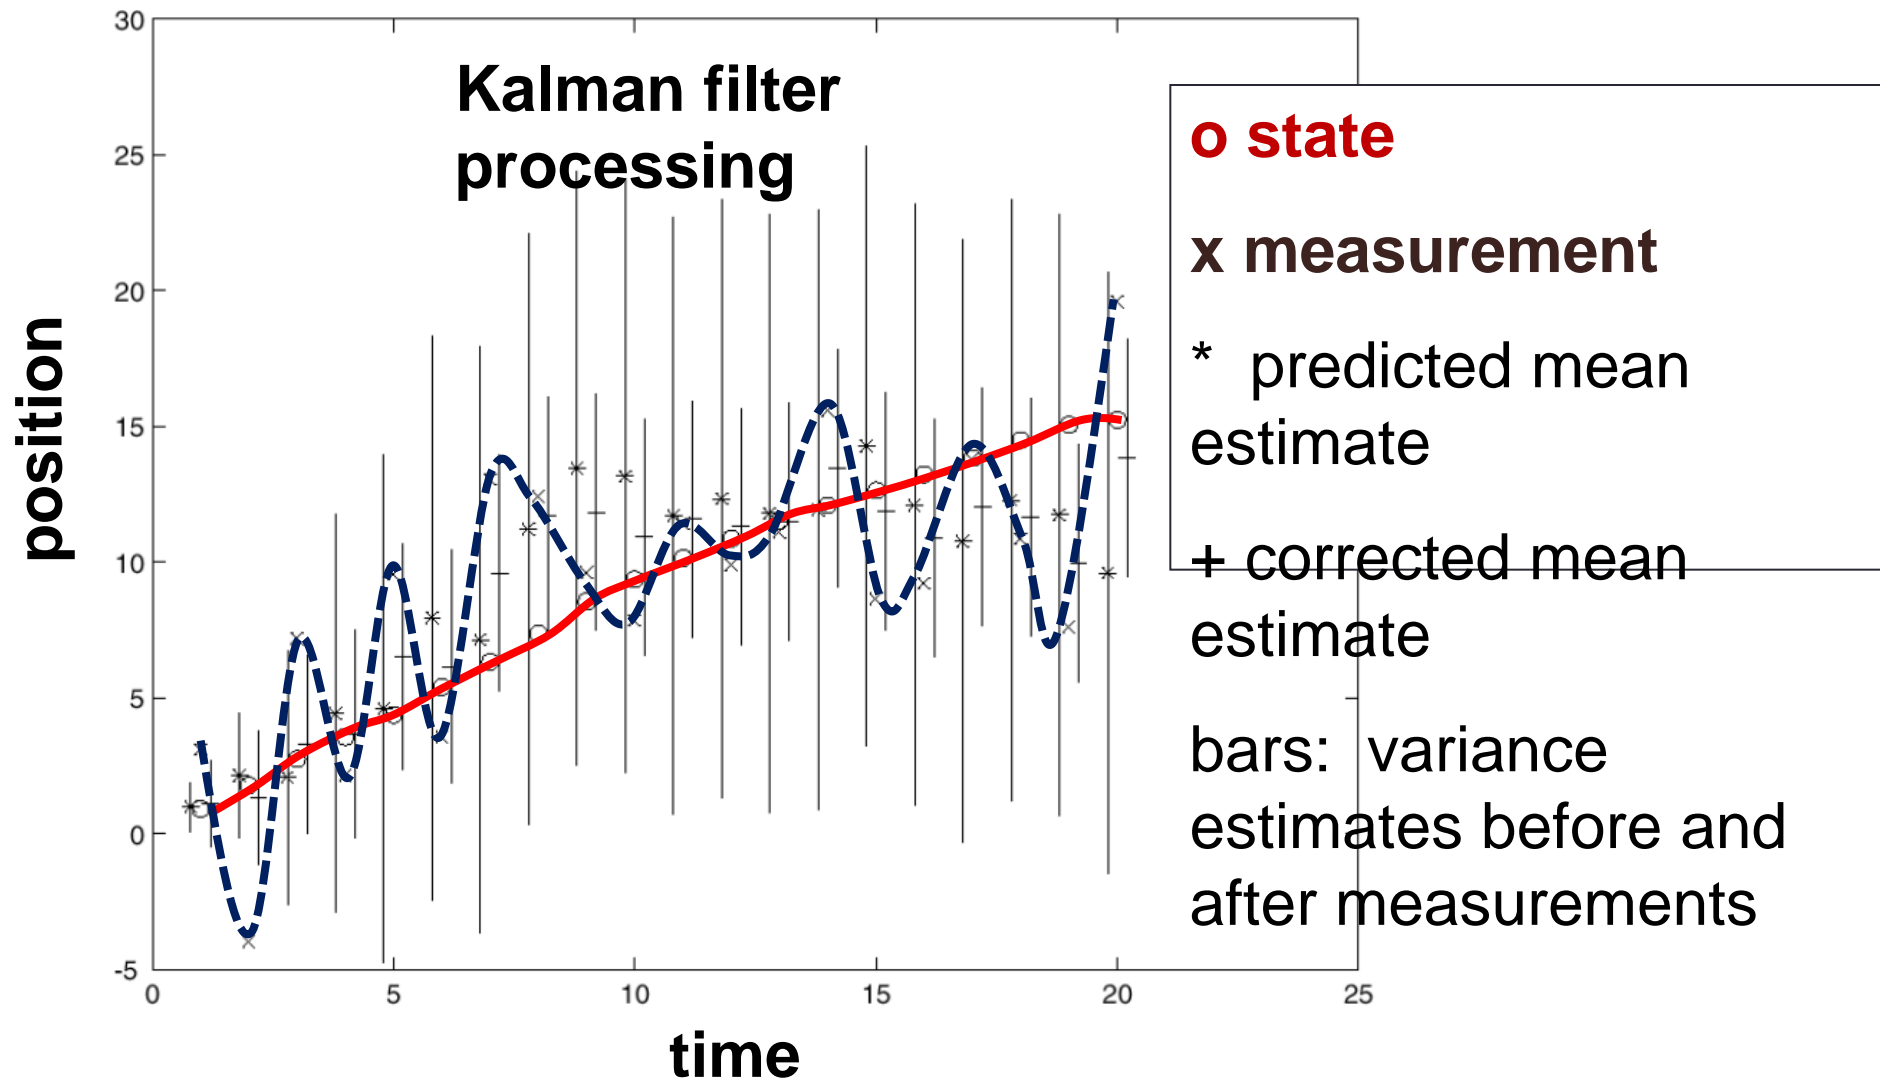
# Constant velocity model



# Constant velocity model




# Constant velocity model





# Constant velocity model



# Kalman filter: General case ( $> 1$ dim)

What if state vectors have more than one dimension?

**PREDICT**

$$x_t^- = D_t x_{t-1}^+$$

$$\Sigma_t^- = D_t \Sigma_{t-1}^+ D_t^T + \Sigma_{d_t}$$

**CORRECT**

$$K_t = \Sigma_t^- M_t^T (M_t \Sigma_t^- M_t^T + \Sigma_{m_t})^{-1}$$

$$x_t^+ = x_t^- + K_t (y_t - M_t x_t^-)$$

$$\Sigma_t^+ = (I - K_t M_t) \Sigma_t^-$$

More weight on residual when measurement error covariance approaches 0.

# Kalman filter pros and cons

- Pros
  - Simple updates, compact and efficient
- Cons
  - Unimodal distribution, only single hypothesis
  - Restricted class of motions defined by linear model
    - Extensions call “Extended Kalman Filtering”
- So what might we do if not Gaussian? Or even unimodal?

*Find out next time! (Actually next Thurs)*