


# CS 4495 Computer Vision

## *Tracking 2: Particle Filters*

Aaron Bobick  
School of Interactive  
Computing



# Administrivia

- PS5 due on Sunday Nov 10, 11:55pm
- PS6 out Thurs Nov 14, due Nov 24<sup>th</sup>
- EXAM: Tues before Thanksgiving, Nov 26<sup>th</sup>    Covers concepts and basics
- Problem set resubmission policy:
  - Full questions only
  - You get 50% credit to replace whatever you got last time on that question.
  - Must be submitted by: DEC 1. NO EXCEPTIONS.


# First some Matlab flow...

# Tracking

- Slides still “adapted” from Kristen Grauman, Deva Ramanan, but mostly from Svetlana Lazebnik
- And now more adaptations from Sebastian Thrun, Dieter Fox and someone who did great particle filter illustrations but whose name is lost to web thievery....

# Recall: Some examples

- <http://www.youtube.com/watch?v=InqV34BcheM>



# Detection vs. tracking



t=1



t=2

...



t=20




t=21

# Detection vs. tracking



Detection: We detect the object independently in each frame and can record its position over time, e.g., based on blob's centroid or detection window coordinates

# Detection vs. tracking



Tracking with *dynamics*: We use image measurements to estimate position of object, but also incorporate position predicted by dynamics, i.e., our expectation of object's motion pattern.




# Tracking with dynamics

- Use model of expected motion to predict where objects will occur in next frame, even before seeing the image.
- **Intent:**
  - Do less work looking for the object, restrict the search.
  - Get improved estimates since measurement noise is tempered by smoothness, dynamics priors.
- **Assumption:** continuous motion patterns:
  - Camera is not moving instantly to new viewpoint
  - Objects do not disappear and reappear in different places in the scene
  - Gradual change in pose between camera and scene

# Tracking as induction

- Base case:
  - Assume we have initial prior that predicts state in absence of any evidence:  $P(X_0)$
  - At the first frame, *correct* this given the value of  $Y_0=y_0$
- Given corrected estimate for frame  $t$ :
  - Predict for frame  $t+1$
  - Correct for frame  $t+1$



# Steps of tracking

- Prediction: What is the next state of the object given past measurements?

$$P(X_t | Y_0 = y_0, \dots, Y_{t-1} = y_{t-1})$$

- Correction: Compute an updated estimate of the state from prediction and measurements

$$P(X_t | Y_0 = y_0, \dots, Y_{t-1} = y_{t-1}, Y_t = y_t)$$

- Tracking can be seen as the process of propagating the **posterior** distribution of state given measurements across time

# Simplifying assumptions

- Only the immediate past matters

$$P(X_t | X_0, \dots, X_{t-1}) = P(X_t | X_{t-1})$$

dynamics model

- Measurements depend only on the current state

$$P(Y_t | X_0, Y_0, \dots, X_{t-1}, Y_{t-1}, X_t) = P(Y_t | X_t)$$

observation model

# Linear Dynamic Models

- Dynamics model: state undergoes linear transformation plus Gaussian noise

- $$X_t \sim N(D_t x_{t-1}, \Sigma_{d_t})$$

- Observation model: measurement is linearly transformed state plus Gaussian noise

$$Y_t \sim N(M_t x_t, \Sigma_{m_t})$$

# Example: Constant velocity (1D)

- State vector is position and velocity

$$x_t = \begin{bmatrix} p_t \\ v_t \end{bmatrix} \quad \begin{aligned} p_t &= p_{t-1} + (\Delta t)v_{t-1} + \varepsilon \\ v_t &= v_{t-1} + \xi \end{aligned} \quad \begin{array}{l} \text{(greek letters} \\ \text{denote noise} \\ \text{terms)} \end{array}$$

$$x_t = D_t x_{t-1} + \text{noise} = \begin{bmatrix} 1 & \Delta t \\ 0 & 1 \end{bmatrix} \begin{bmatrix} p_{t-1} \\ v_{t-1} \end{bmatrix} + \text{noise}$$

- Measurement is position only

$$y_t = Mx_t + \text{noise} = \begin{bmatrix} 1 & 0 \end{bmatrix} \begin{bmatrix} p_t \\ v_t \end{bmatrix} + \text{noise}$$

# Example: Constant acceleration (1D)

- State vector is position, velocity, and acceleration

$$x_t = \begin{bmatrix} p_t \\ v_t \\ a_t \end{bmatrix} \quad \begin{aligned} p_t &= p_{t-1} + (\Delta t)v_{t-1} + \varepsilon \\ v_t &= v_{t-1} + (\Delta t)a_{t-1} + \xi \\ a_t &= a_{t-1} + \zeta \end{aligned} \quad \begin{array}{l} \text{(greek letters} \\ \text{denote noise} \\ \text{terms)} \end{array}$$

$$x_t = D_t x_{t-1} + \text{noise} = \begin{bmatrix} 1 & \Delta t & 0 \\ 0 & 1 & \Delta t \\ 0 & 0 & 1 \end{bmatrix} \begin{bmatrix} p_{t-1} \\ v_{t-1} \\ a_{t-1} \end{bmatrix} + \text{noise}$$

- Measurement is position only


$$y_t = M x_t + \text{noise} = \begin{bmatrix} 1 & 0 & 0 \end{bmatrix} \begin{bmatrix} p_t \\ v_t \\ a_t \end{bmatrix} + \text{noise}$$

# The Kalman filter

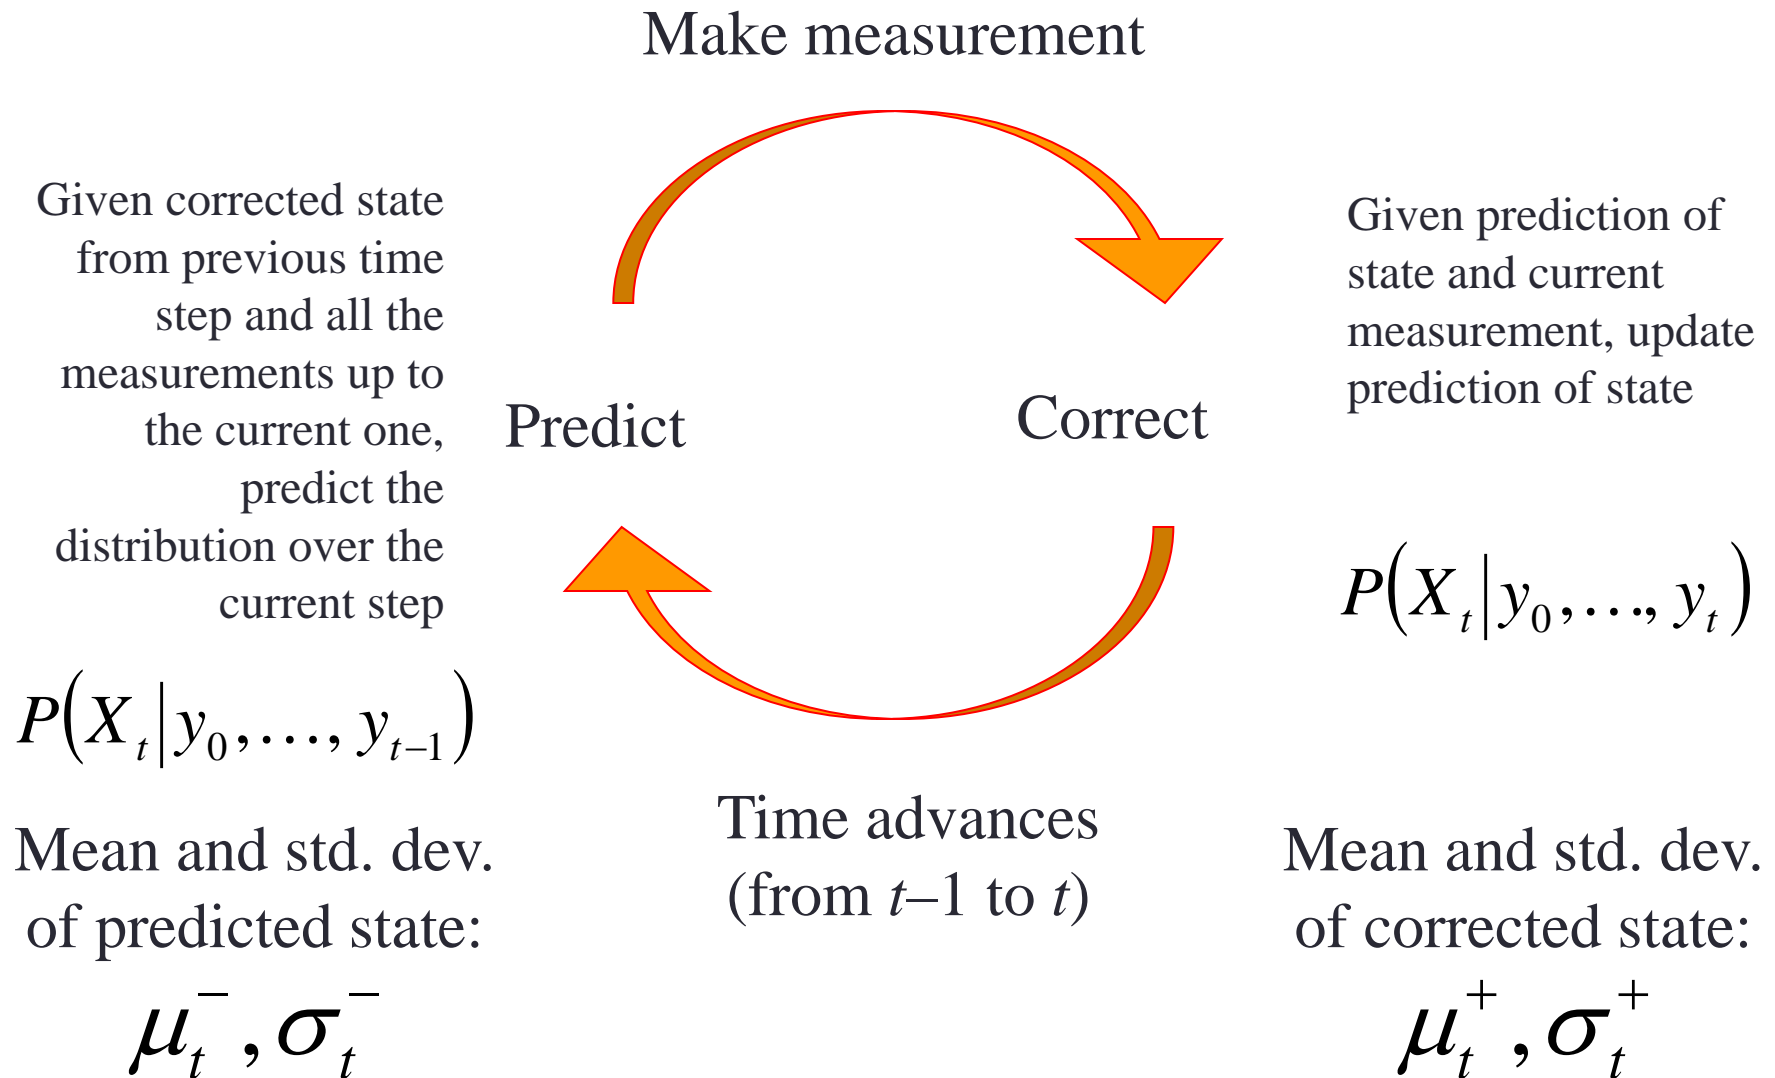
- Method for tracking linear dynamical models in Gaussian noise
- The predicted/corrected state distributions are Gaussian
  - You only need to maintain the mean and covariance
  - The calculations are easy (all the integrals can be done in closed form)



# Propagation of Gaussian densities




# The Kalman Filter: 1D state




# Tracking with KFs: Gaussians!


initial estimate




prediction





measurement



update



# Kalman Filters



$$p(x') = \int p(x' | x) p(x) dx$$

# 1D Kalman filter: Prediction

- Linear dynamic model defines predicted state evolution, with noise

$$X_t \sim N(dx_{t-1}, \sigma_d^2)$$

- Want to estimate distribution for next predicted state

$$P(X_t | y_0, \dots, y_{t-1}) = N(\mu_t^-, (\sigma_t^-)^2)$$

- Update the mean:  $\mu_t^- = d\mu_{t-1}^+$

- Update the variance:  $(\sigma_t^-)^2 = \sigma_d^2 + (d\sigma_{t-1}^+)^2$

# 1D Kalman filter: Correction

- Mapping of state to measurements:  $Y_t \sim N(mx_t, \sigma_m^2)$

- Predicted state:  $P(X_t | y_0, \dots, y_{t-1}) = N(\mu_t^-, (\sigma_t^-)^2)$

- Want to estimate corrected distribution

$$P(X_t | y_0, \dots, y_t) = N(\mu_t^+, (\sigma_t^+)^2)$$

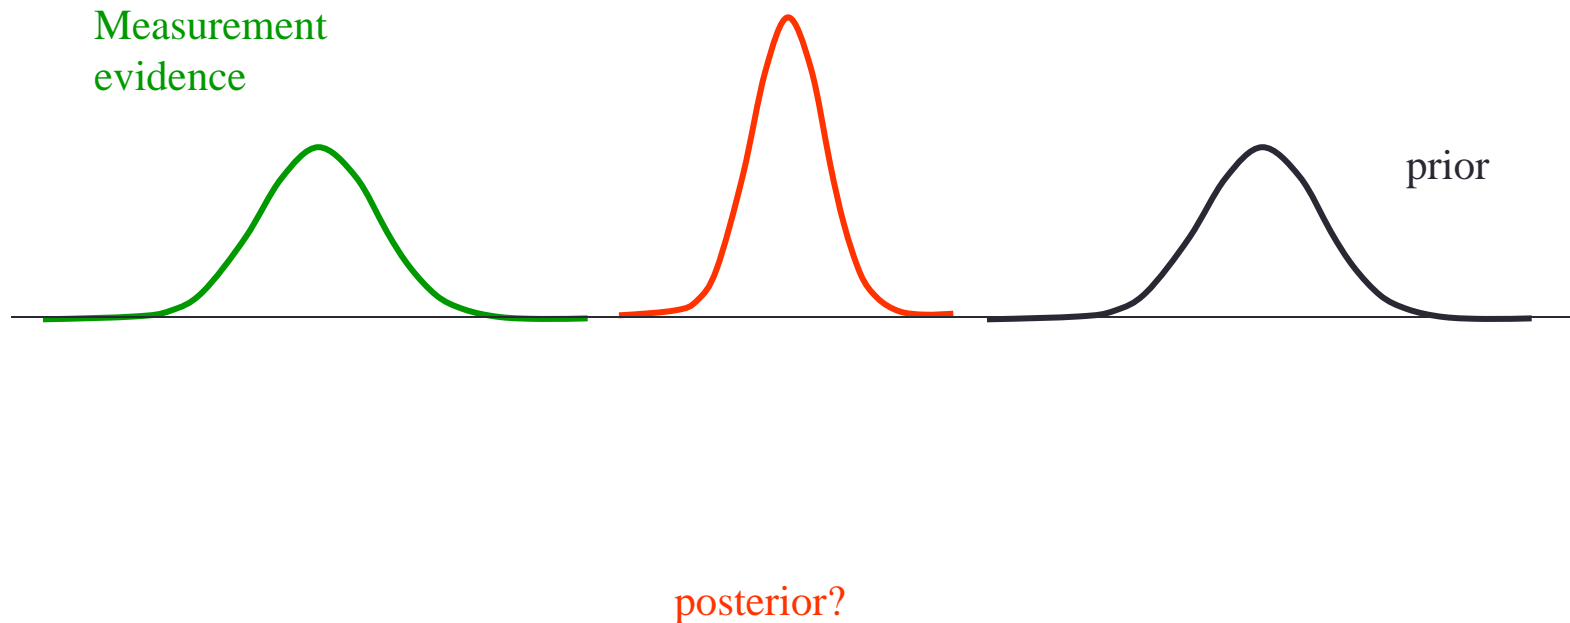
- Update the mean: 
$$\mu_t^+ = \frac{\mu_t^- \sigma_m^2 + m y_t (\sigma_t^-)^2}{\sigma_m^2 + m^2 (\sigma_t^-)^2}$$

- Update the variance:

$$(\sigma_t^+)^2 = \frac{\sigma_m^2 (\sigma_t^-)^2}{\sigma_m^2 + m^2 (\sigma_t^-)^2}$$

# A Quiz

$$p(x | y) \propto p(y | x) p(x)$$




*Does this agree with your intuition?*

# Kalman filter pros and cons

- Pros
  - Simple updates, compact and efficient
- Cons
  - Unimodal distribution, only single hypothesis
  - Restricted class of motions defined by linear model
- So what might we do if not Gaussian?
  - E.g. Object seems to be detected (noisy) in two locations in the next image that are both plausible in terms of dynamics and belief where the object was in the last image.




# Propagation of general densities



# Before we go any further...

- In particle filtering the measurements are written as  $z_t$  and not as  $y_t$
- So we'll start seeing  $z$ 's

# Particle Filters: Basic Idea



→ set of  $n$  (weighted) particles  $X_t$

Density is represented by both **where** the particles are and their **weight**.  $p(x = x_0)$  is now probability of drawing an  $x$  with value  $x_0$ .

Goal:  $p(x_t \in X_t) \approx p(x_t | z_{\{1 \dots t\}})$  with equality when  $n \rightarrow \infty$

# A slight shift...

- For just *tracking* we had the notion of a dynamics model – and for the Kalman filter a *linear* dynamics model.
  - At each time step the prediction was based upon linear transform of state so could easily include velocity, acceleration, etc.
  - The linear form was necessary to make the math work out. To do non-linear dynamics need to linearize about the current state (remember Jacobians?).
- In general, the state transition matrix represented what was known or expected to change from  $t$  to  $t+1$ .
- We can generalize that notion of action or input  $u_t$  at time  $t$ .

# Bayes Filters: Framework

- **Given:**

- Stream of observations  $z$  and action data  $u$ :

$$data_t = \{u_1, z_2 \dots, u_{t-1}, z_t\}$$


- **Sensor model**  $p(z|x)$  .
- **Action model**  $p(x_t|u_{\{t-1\}}, x_{\{t-1\}})$
- **Prior** probability of the system state  $p(x)$ .

- **Wanted:**

- Estimate of the state  $X$  of a **dynamical system**.
- The posterior of the state is also called **Belief**:

$$Bel(x_t) = P(x_t | u_1, z_2 \dots, u_{t-1}, z_t)$$

# Graphical Model Representation



$$p(z_t \mid x_{0:t}, z_{1:t-1}, u_{1:t}) = p(z_t \mid x_t)$$

$$p(x_t \mid x_{1:t-1}, z_{1:t-1}, u_{1:t}) = p(x_t \mid x_{t-1}, u_t)$$

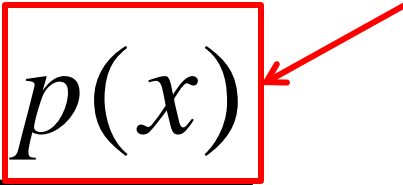
## Underlying Assumptions

- Static world
- Independent noise

# Bayes Rule reminder

$$p(x | z) = \frac{p(z | x) p(x)}{p(z)}$$

Prior before measurement


$$= \eta p(z | x) p(x)$$

# Bayes Filters

$z$  = observation  
 $u$  = action  
 $x$  = state

$$Bel(x_t) = P(x_t \mid u_1, z_2, \dots, u_{t-1}, z_t)$$

**Bayes**  $= \eta P(z_t \mid x_t, u_1, z_2, \dots, u_{t-1}) P(x_t \mid u_1, z_2, \dots, u_{t-1})$

**Markov**  $= \eta P(z_t \mid x_t) P(x_t \mid u_1, z_2, \dots, u_{t-1})$

**Total Probability of “Prior”**  $= \eta P(z_t \mid x_t) \int P(x_t \mid u_1, z_2, \dots, u_{t-1}, x_{t-1}) \cdot P(x_{t-1} \mid u_1, z_2, \dots, u_{t-1}) dx_{t-1}$

**Markov**  $= \eta P(z_t \mid x_t) \int P(x_t \mid u_{t-1}, x_{t-1}) P(x_{t-1} \mid u_1, z_2, \dots, u_{t-1}) dx_{t-1}$

$$= \eta P(z_t \mid x_t) \int P(x_t \mid u_{t-1}, x_{t-1}) Bel(x_{t-1}) dx_{t-1}$$

Prediction *before* taking measurement




# To illustrate...


- Assume a simple robot with a noisy range sensor looking to its side...

# Particle Filters

When density is not easily represented analytically (ie a couple of Gaussians), represent by a set of (possibly weighted) *samples*

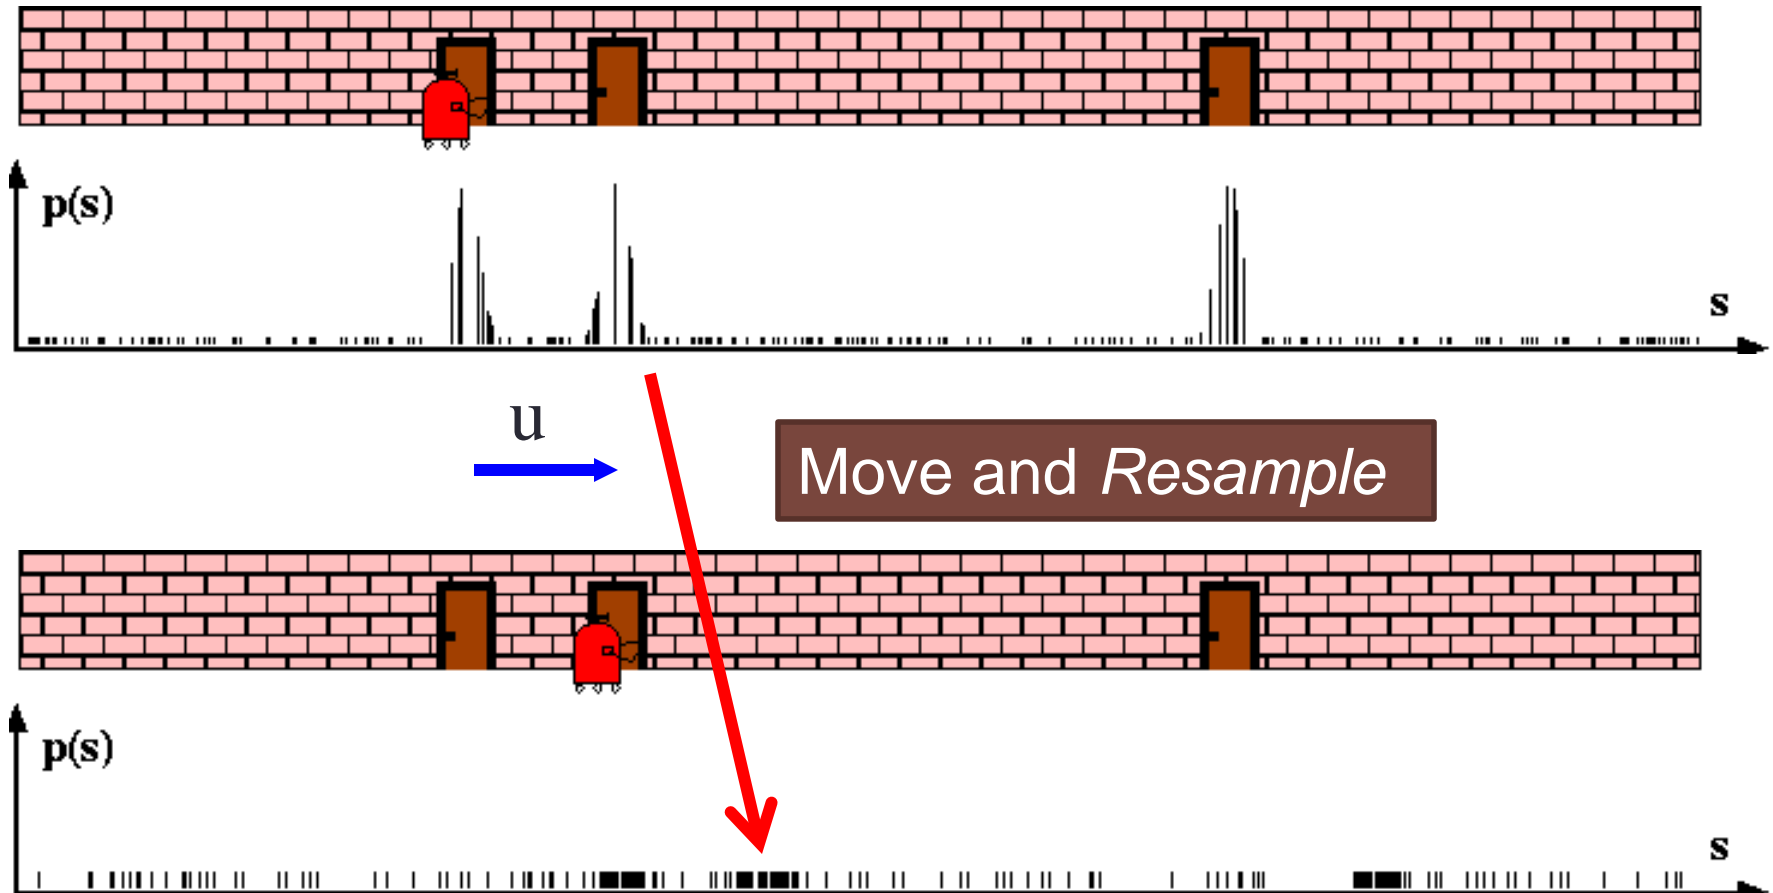


## Sensor Information (aka Importance Sampling)




# Robot Motion


$$Bel^-(x) \leftarrow \int p(x | u, x') Bel(x') dx'$$



## Next Sensor Reading




# Robot Moves Again



# Particle Filter Algorithm (Sequential Importance Resampling)

1. **Algorithm** `particle_filter`  $\{S_{t-1} = \langle x_{t-1}^j, w_{t-1}^j \rangle, u_t, z_t\}$
2.  $S_t = \emptyset, \quad \eta = 0$
3. **For**  $i = 1 \dots n$  *Resample (generate  $i$  new samples)*
4.     Sample index  $j(i)$  from the discrete distribution given by  $w_{t-1}$
5.     Sample  $x_t^i$  from  $p(x_t | x_{t-1}, u_t)$  using  $x_{t-1}^{j(i)}$  and  $u_t$  *Control*
6.      $w_t^i = p(z_t | x_t^i)$  *Compute importance weight (or reweight)*
7.      $\eta = \eta + w_t^i$  *Update normalization factor*
8.      $S_t = S_t \cup \{\langle x_t^i, w_t^i \rangle\}$  *Insert*
9. **For**  $i = 1 \dots n$
10.      $w_t^i = w_t^i / \eta$  *Normalize weights*

# Graphical steps particle filtering

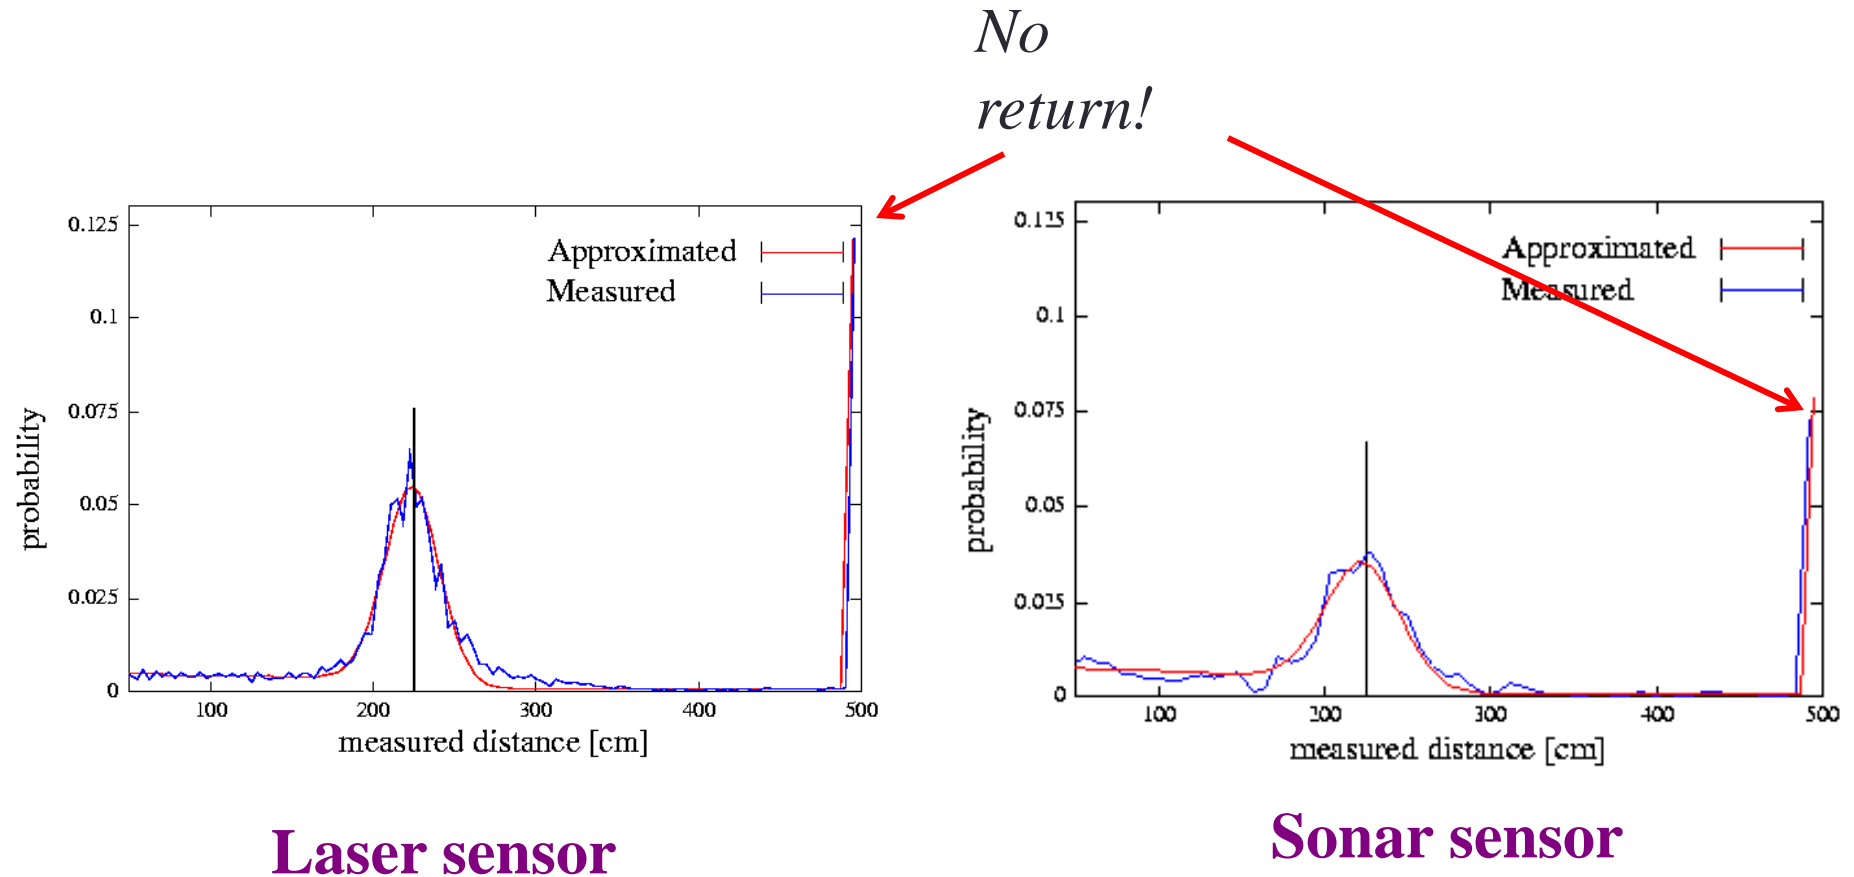





# A real robot sensing problem

- Assume a robot **knows** a 3D map of its world
- It has a noisy depth sensor or sensors whose sensing uncertainty is **known**.
- It moves from frame to frame.
- How well can it know where it is  $(x, y, \theta)$  ?


# Proximity Sensor Model



# Sonar Example: Initial Distribution




# After Incorporating Ten Ultrasound Scans




# After Incorporating 65 Ultrasound Scans



# Estimated Path





# Sample-based Localization (sonar)





# *Smithsonian Museum of American History...*





















































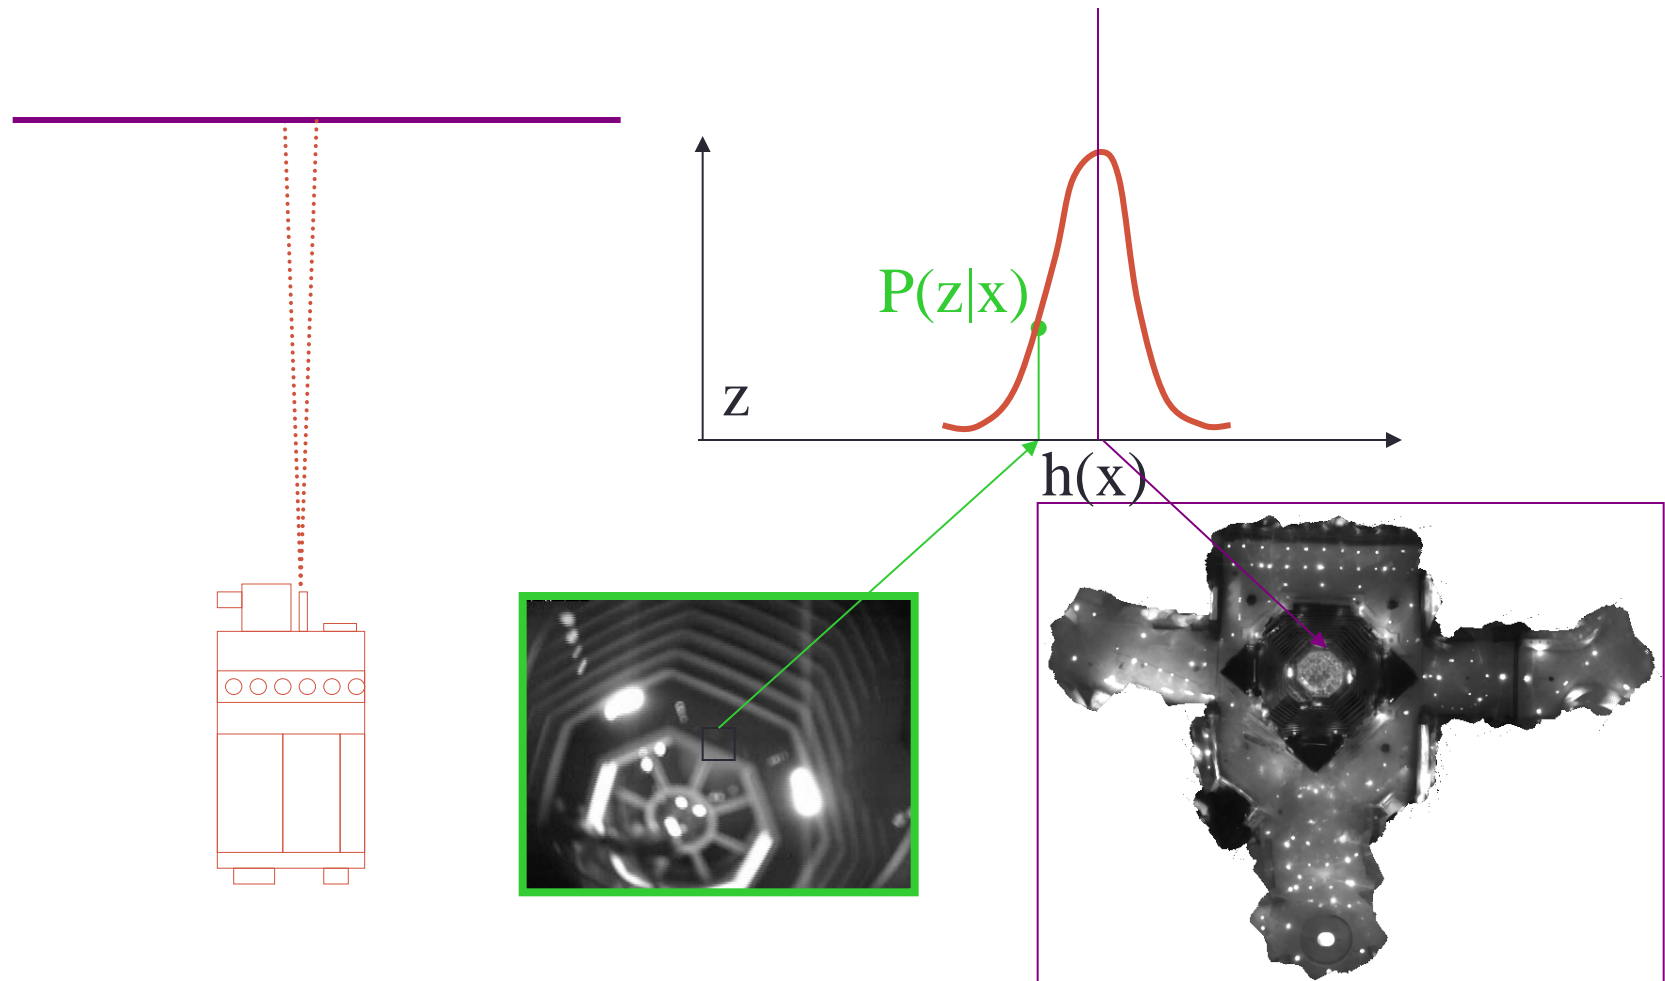
*How about simple vision....*

# Using Ceiling Maps for Localization




Dellaert, et al. 1997

# Vision-based Localization



# Under a Light

**Measurement  $z$ :**




**$P(z/x)$ :**




# Next to a Light

**Measurement  $z$ :**



**$P(z/x)$ :**




# Elsewhere


Measurement  $z$ :



$P(z/x)$ :




# Global Localization Using Vision








# Using Vision




# A detail: Resampling method can matter

- **Given**: Set  $S$  of weighted samples.
- **Wanted** : Random sample, where the probability of drawing  $x_i$  is given by  $w_i$ .
- Typically done  $n$  times with replacement to generate new sample set  $S'$ .
  - *Or even not done except when needed...too many low weight particles.*

# Resampling



- Roulette wheel
- Binary search,  $n \log n$



- Stochastic universal sampling
- Systematic resampling
- Linear time complexity
- Easy to implement, low variance

# Resampling Algorithm

1. Algorithm **systematic\_resampling**( $S, n$ ):
2.  $S' = \emptyset, c_1 = w^1$
3. **For**  $i = 2 \dots n$  *Generate cdf*
4.  $c_i = c_{i-1} + w^i$
5.  $u_1 \sim U[0, n^{-1}]$ ,  $i = 1$  *Initialize threshold*
6. **For**  $j = 1 \dots n$  *Draw samples ...*
7. **While** (  $u_j > c_i$  ) *Skip until next threshold reached*
8.  $i = i + 1$
9.  $S' = S' \cup \{ \langle x^i, n^{-1} \rangle \}$  *Insert*
10.  $u_{j+1} = u_j + n^{-1}$  *Increment threshold*
11. **Return**  $S'$  (Also called *stochastic universal sampling*)

# PF: Practical Considerations

- If dealing with **highly peaked observations**
  - Add noise to observation and prediction models
  - Better proposal distributions: e.g., perform Kalman filter step to determine proposal
- **Overestimating noise** often reduces number of required samples
- **Recover from failure** by selectively adding samples from observations
- **Recover from failure** by uniformly adding some samples
- Can **Resample** only when necessary (efficiency of representation measured by variance of weights)

# Tracking issues

- Next time?

# Tracking issues

- Initialization
  - Manual
  - Background subtraction
  - Detection



# Tracking issues

- Initialization
- Obtaining observation and dynamics model
  - Dynamics model: learn (very difficult) or specify using domain knowledge
  - Generative observation model: “render” the state on top of the image and compare

# Tracking issues


- Initialization
- Obtaining observation and dynamics model
- Prediction vs. correction
  - If the dynamics model is too strong, will end up ignoring the data
  - If the observation model is too strong, tracking is reduced to repeated detection

# Tracking issues

- Initialization
- Obtaining observation and dynamics model
- Prediction vs. correction
- Data association
  - What if we don't know which measurements to associate with which tracks?


# Data association

- So far, we've assumed the entire measurement to be relevant to determining the state
- In reality, there may be uninformative measurements (clutter) or measurements may belong to different tracked objects
- **Data association:** task of determining which measurements go with which tracks




# Data association

- Simple strategy: only pay attention to the measurement that is “closest” to the prediction



# Data association

- Simple strategy: only pay attention to the measurement that is “closest” to the prediction



Doesn't always work...

Alternative: keep track of **multiple hypotheses** at once...

# Data association


- Simple strategy: only pay attention to the measurement that is “closest” to the prediction
- More sophisticated strategy: keep track of multiple state/observation hypotheses
  - Can be done with particle filtering
- This is a general problem in computer vision, there is no easy solution

# Tracking issues

- Initialization
- Obtaining observation and dynamics model
- Prediction vs. correction
- Data association
- Drift
  - Errors caused by dynamical model, observation model, and data association tend to accumulate over time



# Drift



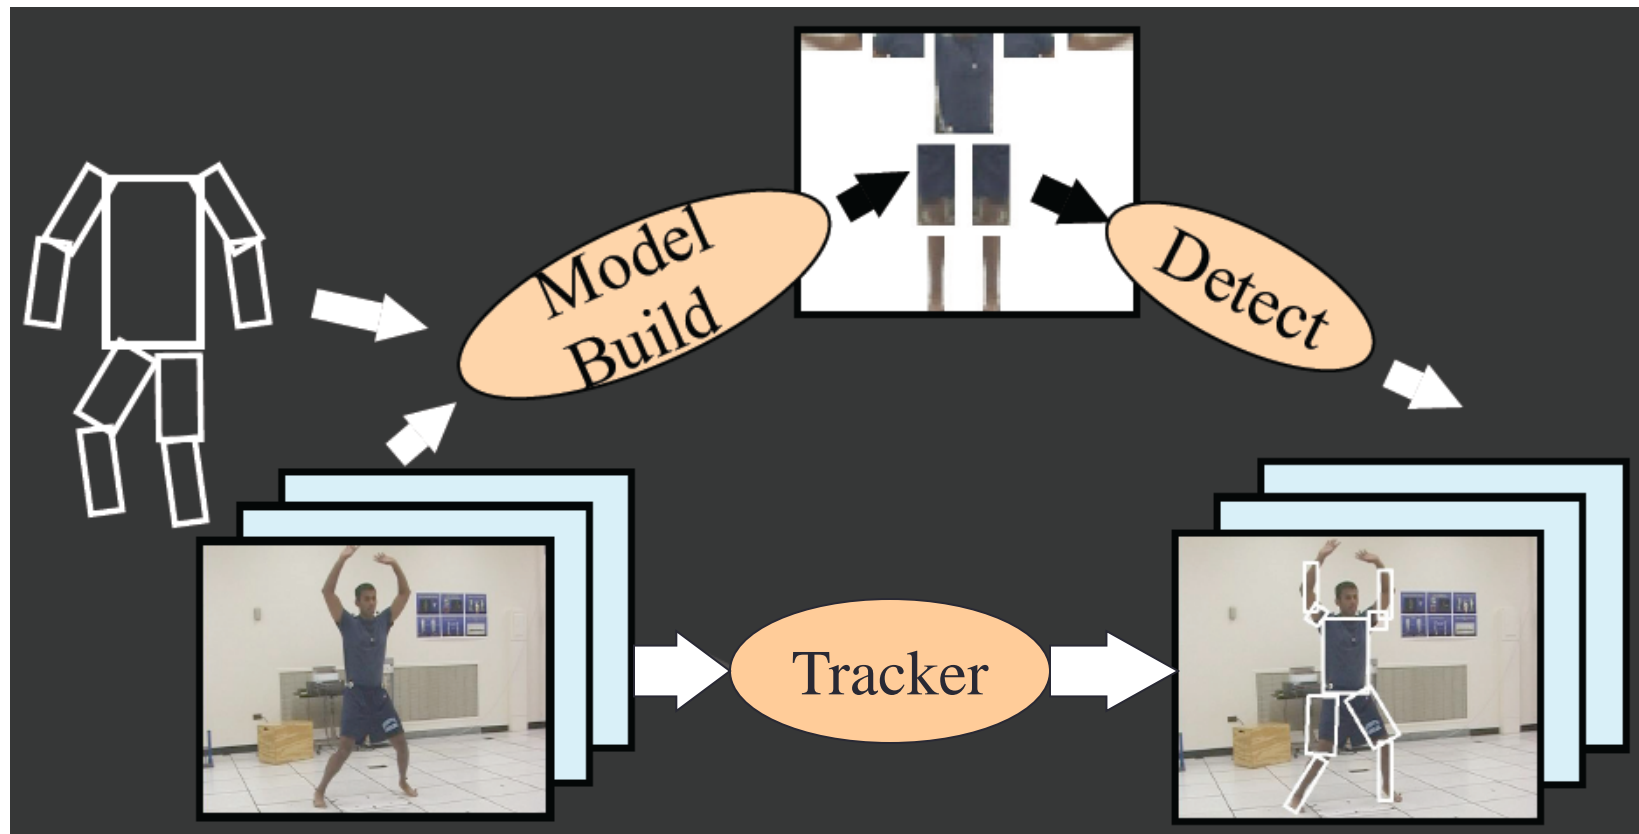
D. Ramanan, D. Forsyth, and A. Zisserman. [Tracking People by Learning their Appearance](#). PAMI 2007.

# Tracking people by learning their appearance

- Person model = appearance + structure (+ dynamics)
- Structure and dynamics are generic, appearance is person-specific
- Trying to acquire an appearance model “on the fly” can lead to drift
- Instead, can use the whole sequence to initialize the appearance model and then keep it fixed while tracking
- Given strong structure and appearance models, tracking can essentially be done by repeated detection (with some smoothing)

D. Ramanan, D. Forsyth, and A. Zisserman. [Tracking People by Learning their Appearance](#). PAMI 2007.




# Tracking people by learning their appearance



D. Ramanan, D. Forsyth, and A. Zisserman. [Tracking People by Learning their Appearance](#). PAMI 2007.

# Pictorial structure model

Fischler and Elschlager(73), Felzenszwalb and Huttenlocher(00)




$$\Pr(P_{\text{tor}}, P_{\text{arm}}, \dots | \text{Im}) \propto \prod_{i,j} \Pr(P_i | P_j) \prod_i \Pr(\text{Im}(P_i))$$

↑
↑

part geometry
part appearance




# Bottom-up initialization: Clustering



D. Ramanan, D. Forsyth, and A. Zisserman. [Tracking People by Learning their Appearance](#). PAMI 2007.


# Top-down initialization: Exploit “easy” poses



D. Ramanan, D. Forsyth, and A. Zisserman. [Tracking People by Learning their Appearance](#). PAMI 2007.




# Tracking by model detection



D. Ramanan, D. Forsyth, and A. Zisserman. [Tracking People by Learning their Appearance](#). PAMI 2007.

# Example results



<http://www.ics.uci.edu/~dramanan/papers/pose/index.html>